


# FLOATING POINT NUMBERS



IEEE floating point standard

---


**Bernhard Kainz** (with thanks to **A. Gopalan**, **N. Dulay** and **E. Edwards**)

[b.kainz@imperial.ac.uk](mailto:b.kainz@imperial.ac.uk)

# IEEE floating point standard

- **IEEE: institute of electrical and electronic engineers (USA)**
- **Comprehensive standard** for binary floating point arithmetic
- Widely adopted → **predictable results** independent of architecture
- Standard defines:
  - **Format** of binary floating point numbers, i.e. how the fields are stored in memory
  - **Semantics** of arithmetic operations
  - Rules for **error conditions**

# Single precision format (32-bit)



- Coefficient is called the **significand** in the IEEE standard
- Value represented is  $\pm 1.F \times 2^{E-127}$
- The **normal bit** (the 1.) is omitted from the significand field → a **hidden bit**
- Single precision yields **24 bits** (approx. **7 decimal digits** of precision)
- **Normalised ranges** in decimal are approximately:  
 $-10^{38}$  to  $-10^{-38}$ ,    **0**,     $10^{38}$  to  $10^{-38}$

# Exponent field


- In the IEEE standard, exponents are stored as **excess values**, not as 2's complement

- Example: **In 8-bit excess-127**

-127	would be held as	0000 0000
...		...
0		0111 1111
1		1000 0000
...		...
128		1111 1111

- Allows non-negative floating point numbers to be compared using simple integer comparisons

# Double precision format (64-bit)



- Value represented is  $\pm 1.F \times 2^{E-1023}$
- Double precision yields **53 bits** (approx. **16 decimal digits** of precision)
- **Normalised ranges** in decimal are approximately:  
 $-10^{308}$  to  $-10^{-308}$ , **0**,  $10^{308}$  to  $10^{-308}$
- Single precision generally reserved for when memory is scarce or for debugging numerical calculations since rounding errors show up more quickly

# Example: conversion to IEEE format

What is 42.6875 in IEEE single precision format?

1. Convert to **binary number**:  $42.6875 = 10\ 1010.1011$
2. **Normalise**:  $1.0101\ 0101\ 1 \times 2^5$
3. **Significand field** is thus:  
 $0101\ 0101\ 1000\ 0000\ 0000\ 0000$
4. **Exponent field** is  $(5 + 127 = 132)$ :  $1000\ 0100$

Sign S	Exponent E	Significand F
0	1000 0100	0101 0101 1000 0000 0000 0000

Hex: 422A C000

# Example: conversion from IEEE format

What is the IEEE single precision value represented by **BEC0 0000** in decimal?

Sign S	Exponent E	Significand F
1	0111 1101	1000 0000 0000 0000 0000 000

1. **Exponent field:**  $0111\ 1101 = 125$
2. **True binary exponent:**  $125 - 127 = -2$
3. **Significand field + hidden bit:**  
 $1.1000\ 0000\ 0000\ 0000\ 0000\ 000$
4. So **unsigned value** is  $1.1 \times 2^{-2} = 0.011$  (binary)  
 $= 0.25 + 0.125 = 0.375$  (decimal)
5. Adding **sign bit** gives finally  $-0.375$

# Example: addition

Carry out the addition  $42.6875 + 0.375$  in IEEE single precision arithmetic

Number	Sign	Exponent	Significand
42.6875	0	1000 0100	0101 0101 1000 0000 0000 000
0.375	0	0111 1101	1000 0000 0000 0000 0000 000

- To add these numbers, exponents must be the same → make the smaller exponent equal to the larger by shifting significand accordingly
- **Note:** must restore **hidden bit** when carrying out floating point operations



# Example: addition (cont.)

- **Significand** of larger no.: 1.0101 0101 1000 0000 0000 000
- **Significand** of smaller no.: 1.1000 0000 0000 0000 0000 000
- Exponents differ by  $(1000\ 0100 - 0111\ 1101 = 7)$  so shift binary point of smaller no. 7 places to the left:
- **Significand** of smaller no.: 0.0000 0011 0000 0000 0000 000
- **Significand** of larger no.: 1.0101 0101 1000 0000 0000 000
- **Significand** of **sum**: 1.0101 1000 1000 0000 0000 000
- So **sum** is  $1.0101\ 1000\ 1 \times 2^5 = 10\ 1011.0001 = 43.0625$

Sign S	Exponent E	Significand F
0	1000 0100	0101 1000 1000 0000 0000 000

# Special values

- IEEE formats can encode five kinds of values: **zero**, **normalised numbers**, **denormalised numbers**, **infinity** and **not-a-number (NaNs)**
- Single precision representations:

IEEE value	Sign field	Exponent	Significand	True exponent
$\pm 0$	0 or 1	0	0 (all zeros)	
$\pm$ denormalised no.	0 or 1	0	Any non-zero bit pattern	-126
$\pm$ normalised no.	0 or 1	1 ... 254	Any bit pattern	-126 ... 127
$\pm \infty$	0 or 1	255	0 (all zeros)	
Not-a-number	0 or 1	255	Any non-zero bit pattern	

# Denormalised numbers

- An **all zero exponent** is used to represent both **zero** and **denormalised numbers**
- An **all one exponent** is used to represent **infinities** and **not-a-numbers**
- Means **range for normalised numbers is reduced**, for single precision the exponent range is  $-126 \dots 127$  rather than  $-127 \dots 128$
- **Denormalised numbers** represent values between the underflow limits and zero, i.e. for single precision we have  $\pm 0.F \times 2^{-126}$
- Allows a more **gradual shift to zero** – useful in some numerical applications

# Infinites and NaNs

- Infinites represent values **exceeding the overflow limits** and for divisions of non-zero quantities by zero
- You can do basic ‘arithmetic’ with them, e.g.:

$$\infty + 5 = \infty, \quad \infty + \infty = \infty$$

- NaNs represent the result of operations which have **no (real) mathematical interpretation**, e.g.

$$\frac{0}{0}, \quad +\infty + -\infty, \quad 0 \times \infty, \quad \text{square root of a negative number}$$

- Operations resulting in NaNs can either yield a NaN result (**quiet NaN**) or an exception (**signalling NaN**)

# Special Operations

Operation	Result
$N \div \pm \text{Infinity}$	0
$\pm \text{Infinity} \times \pm \text{Infinity}$	$\pm \text{Infinity}$
$\pm \text{non-zero} \div 0$	$\pm \text{Infinity}$
$\text{Infinity} + \text{Infinity}$	Infinity
$\pm 0 \div \pm 0$	<i>NaN</i>
$\text{Infinity} - \text{Infinity}$	<i>NaN</i>
$\pm \text{Infinity} \div \pm \text{Infinity}$	<i>NaN</i>
$\pm \text{Infinity} \times 0$	<i>NaN</i>

😊 SOME FUN 😊

---

# Floating Point Precision

- C code:

```
#include <stdio.h>
int main() {

    float a, b, c;

    float EPSILON = 0.0000001;

    a = 1.345f; b = 1.123f;

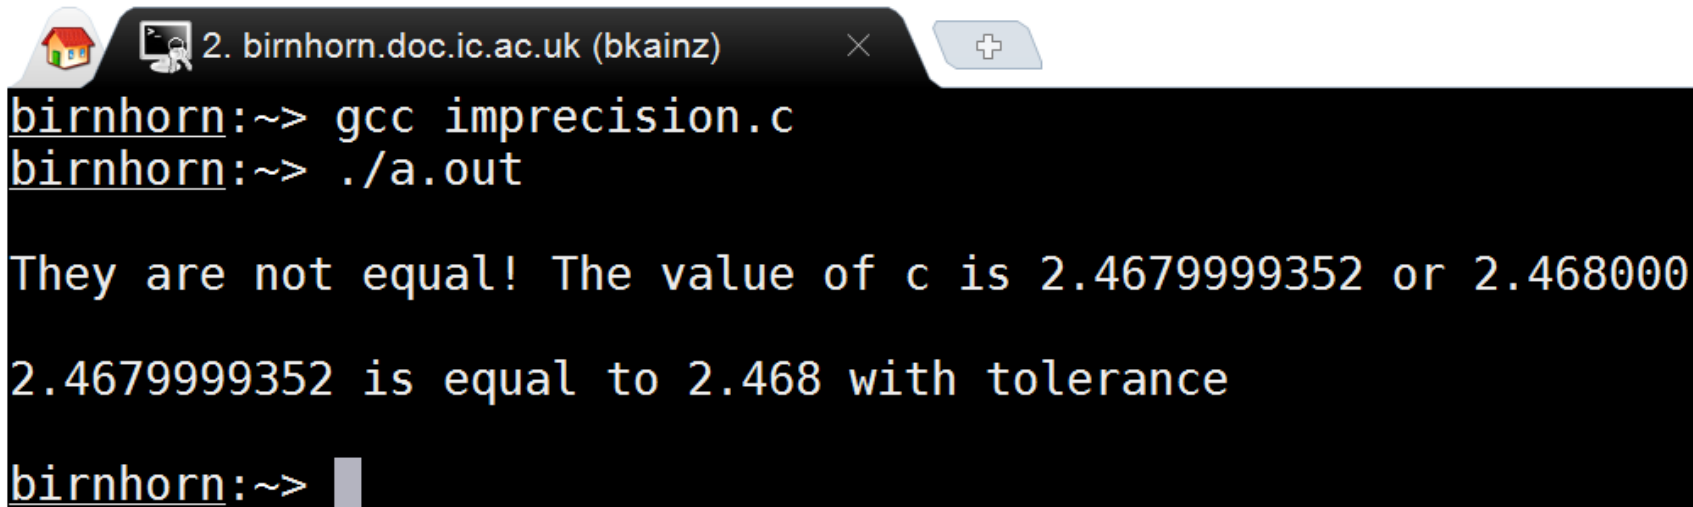
    c = a + b;

    if (c == 2.468)
        printf ("They are equal.\n");
    else
        printf ("\nThey are not equal! The value of c is %.10f or %f\n",c,c);

    // With some tolerance

    if (((2.468 - EPSILON) < c) && (c < (2.468 + EPSILON)))
        printf ("\n%.10f is equal to 2.468 with tolerance\n\n", c);
}
```

# Run-time



A terminal window with a dark background. The title bar shows a home icon, a terminal icon, and the text "2. birnhorn.doc.ic.ac.uk (bkainz)". The terminal content shows the compilation of a C program named "imprecision.c" using "gcc", followed by its execution. The output states that two values are not equal, showing a long decimal representation of a value. A comment explains that the values are equal within a certain tolerance. The prompt "birnhorn:~>" is followed by a cursor.

```
birnhorn:~> gcc imprecision.c
birnhorn:~> ./a.out

They are not equal! The value of c is 2.4679999352 or 2.468000

2.4679999352 is equal to 2.468 with tolerance

birnhorn:~> █
```



# Finding Machine Epsilon

- Pseudo-code

Set machineEps = 1.0;

Loop

    machineEps = machineEps/2.0

Until  $((1 + \text{machineEps}/2.0) \neq 1)$

Print machineEps

# Finding Machine Epsilon

- C code

```
#include <stdio.h>
```

```
int main( int argc, char **argv )
```

```
{
```

```
    float machEps = 1.0f;
```

```
    do {
```

```
        machEps /= 2.0f;
```

```
        // If next epsilon yields 1, then break, because current
```

```
        // epsilon is the machine epsilon.
```

```
    }
```

```
    while ((float)(1.0 + (machEps/2.0f)) != 1.0);
```

```
    printf( "\nCalculated Machine epsilon: %G\n\n", machEps );
```

```
    return 0;
```

```
}
```

# Finding Machine Epsilon

- In Java

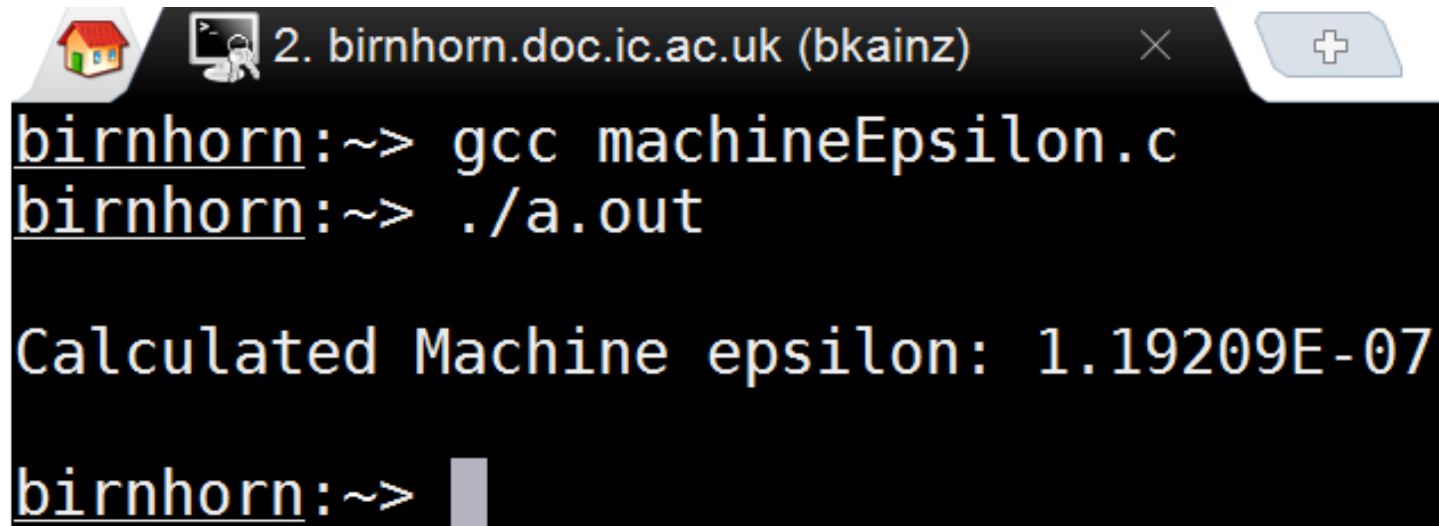
```
public class machEps
{
    private static void calculateMachineEpsilonFloat() {
        float machEps = 1.0f;

        do {
            machEps /= 2.0f;
        } while ((float)(1.0 + (machEps/2.0)) != 1.0);

        System.out.println( "Calculated machine epsilon: " + machEps );
    }

    public static void main (String args[])
    {
        calculateMachineEpsilonFloat ();
    }
}
```

# Run-time



A terminal window with a dark background and light-colored text. The window title bar shows a home icon, a terminal icon, and the text "2. birnhorn.doc.ic.ac.uk (bkainz)". The terminal content shows the user "birnhorn" at the prompt "~>" running the command "gcc machineEpsilon.c", followed by another prompt "~>" where the command "./a.out" is entered. The output of the program is "Calculated Machine epsilon: 1.19209E-07". The final prompt "~>" is followed by a small grey rectangular cursor.

```
2. birnhorn.doc.ic.ac.uk (bkainz)
birnhorn:~> gcc machineEpsilon.c
birnhorn:~> ./a.out

Calculated Machine epsilon: 1.19209E-07

birnhorn:~> █
```

# Special Operations

- Example

```
#include <stdio.h>
```

```
int main (int argc, char **argv)
```

```
{
```

```
    float a = 1.0/0.0;
```

```
    float b = a * -100;
```

```
    float c = b/a;
```

```
    int d = 2 * 10 + 3;
```

```
    printf ("\nValue of a = %f\n\n", a);
```

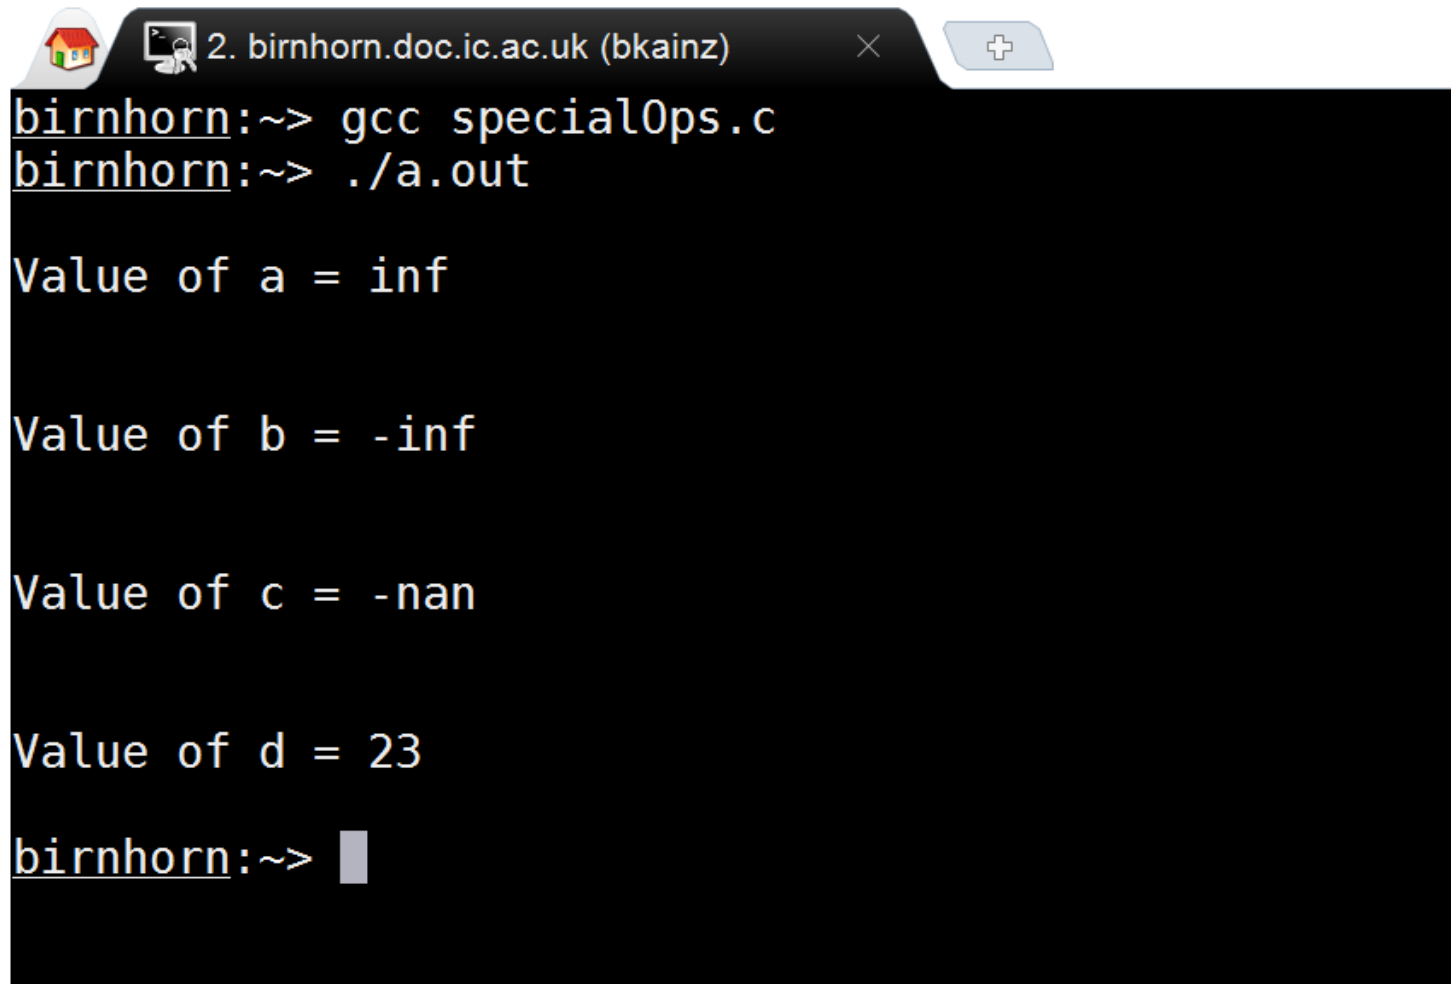
```
    printf ("\nValue of b = %f\n\n", b);
```

```
    printf ("\nValue of c = %f\n\n", c);
```

```
    printf ("\nValue of d = %d\n\n", d);
```

```
}
```

# Run-time



A terminal window with a dark background and light-colored text. The window has a title bar with a home icon, a document icon, and the text "2. birnhorn.doc.ic.ac.uk (bkainz)". There are also close, maximize, and window control buttons. The terminal shows the following commands and output:

```
birnhorn:~> gcc special0ps.c
birnhorn:~> ./a.out

Value of a = inf

Value of b = -inf

Value of c = -nan

Value of d = 23

birnhorn:~> █
```