

# Communication and Synchronisation

Files

Signals (UNIX)

Events, exceptions (Windows)

Pipes

Message Queues (UNIX)

Mailslots (Windows)


Sockets – in NDS course

Shared memory

Semaphores, Locks, Monitors


# Types of Process Interaction

## Sharing




Require **mutually exclusive** access to prevent interference

## Synchronisation



P1 informs P2 that some event has happened. P2 waits for it

## Communication



P1 sends P2 data. P1 blocks when buffer is full, P2 blocks when buffer is empty.

**Mutual Exclusion, Synchronisation and Communication** are closely related.

# UNIX Signals

Inter-Process Communication (IPC) mechanism

Signal delivery similar to delivery of hardware interrupts

Used to notify processes when an event occurs

A process can send a signal to another process if it has permission

- *"the real or effective user ID of the receiving process must match that of the sending process or the user must have appropriate privileges (such as given by a set-user-ID program or the user is the super-user)." (man page)*
- The kernel can send signals to any process

# When Are Signals Generated?

When an exception occurs

- e.g., division by zero => **SIGFPE**,  
segment violation => **SIGSEGV**

When the kernel wants to notify the process of an event

- e.g., if process writes to a closed pipe => **SIGPIPE**

When certain key combinations are typed in a terminal

- e.g., Ctrl-C => **SIGINT**

Explicitly using the **kill()** system call

# UNIX Signals – Examples

<b>SIGINT</b>	Interrupt from keyboard
<b>SIGABRT</b>	Abort signal from <b>abort</b>
<b>SIGFPE</b>	Floating point exception
<b>SIGKILL</b>	Kill signal
<b>SIGSEGV</b>	Invalid memory reference
<b>SIGPIPE</b>	Broken pipe: write to pipe with no readers
<b>SIGALRM</b>	Timer signal from <b>alarm</b>
<b>SIGTERM</b>	Termination signal

# UNIX Signals

The default action for most signals is to terminate the process

But the receiving process may choose to

- Ignore it
- Handle it by installing a signal handler
- Two signals cannot be ignored/handled: **SIGKILL** and **SIGSTOP**

```
signal(SIGINT, my_handler);  
  
void my_handler(int sig) {  
    printf("Received SIGINT. Ignoring...")  
}
```

# Signal Handlers – Example

```
#include <signal.h>
#include <stdio.h>

void my_handler(int sig) {
    fprintf(stderr, "SIGINT caught!");
}

int main(int argc, char *argv[])
{
    signal(SIGINT, my_handler);
    while (1) {}
}
```

```
$ ./a.out
[ctrl-C]
SIGINT caught
```

# Sockets

Allow bidirectional communication

Can be used to exchange information both locally and across a network

- Unlike pipes which are identified by machine specific file descriptors

Two types of sockets:

- TCP (stream sockets)
- UDP (datagram sockets)

Covered in Networks and Distributed Systems course




# Shared Memory

Processes can set up shared memory areas

- Implicitly or explicitly mapped to files on disk

After shared memory is established, no need for kernel involvement



# Shared Memory – System V API

<b>shmget</b>	Allocates a shared memory segment
<b>shmat</b>	Attaches a shared memory segment to the address space of a process
<b>shmctl</b>	Changes the properties associate with a shared memory segment
<b>shmdt</b>	Detaches a shared memory segment from a process

# Synchronisation

# Process Synchronization

How do processes synchronize their operation to perform a task?

Key concepts:

- Critical sections
- Mutual exclusion
- Atomic operations
- Race conditions
- Synchronization mechanisms
  - ✦ Locks, semaphores, monitors, etc.
- Deadlock
- Starvation


Concepts relevant to both **processes** and **threads**

# Shared Data Example

Account #1234: £10,000



Extract £1000  
from account 1234



Extract £1000  
from account 1234

# Shared Data Example

```
void Extract(int acc_no, int sum)
{
    int B = Acc[acc_no];
    Acc[acc_no] = B - sum;
}
```

Acc[1234] 10,000

B = 10,000  
Acc[1234] = 9000



Extract(1234, 1000)

B = 9,000  
Acc[1234] = 8000



Extract(1234, 1000)

# Shared Data Example

```
void Extract(int acc_no, int sum)
{
    int B = Acc[acc_no];
    Acc[acc_no] = B - sum;
}
```

**Critical section!**  
Need **mutual exclusion**

Acc[1234] 10,000

B = 10,000

Acc[1234] = 9000

B = 10,000

Acc[1234] = 9000



Extract(1234, 1000)

Extract(1234, 1000)

# Critical Sections and Mutual Exclusion

***Critical section/region***: section of code in which processes access a shared resource – executed by only one process at a time.

A code section is critical if it:

1. Reads a memory location which is shared with another process
2. Updates a shared memory location with a value which depends on what it read

***Mutual exclusion*** ensures that if a process is executing its critical section, no other process can be executing it

- Processes must request ***permission*** to enter critical sections

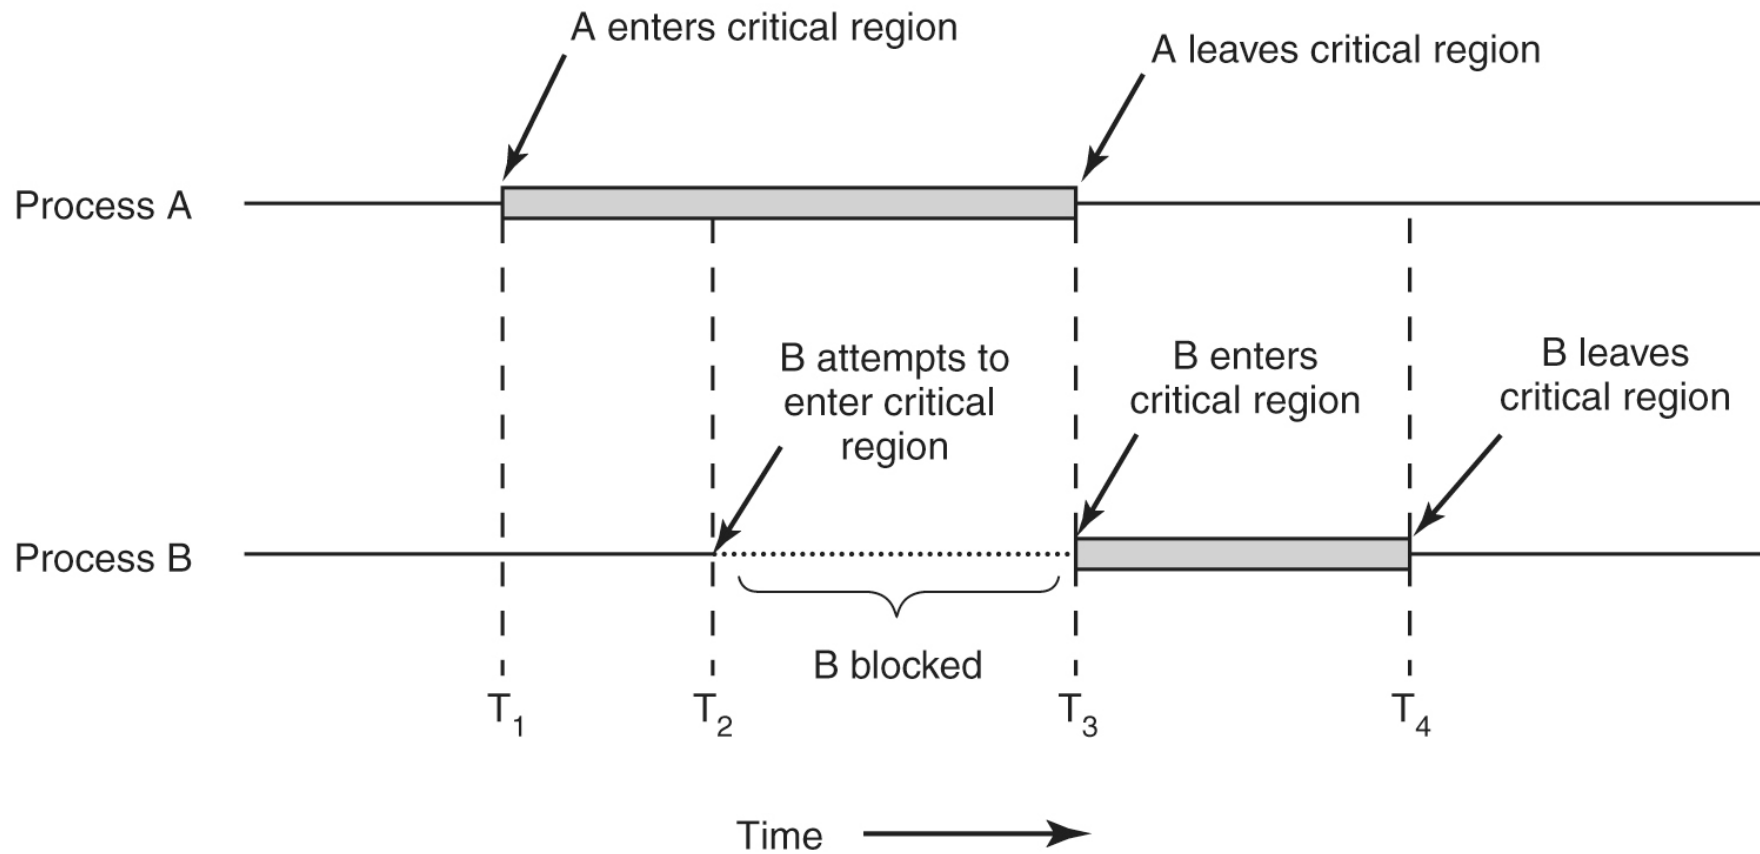
A ***synchronisation mechanism*** is required at the entry and exit of the critical section



# Requirements for Mutual Exclusion

- No two processes may be simultaneously inside a critical section
- No process running outside the critical section may prevent other processes from entering the critical section
  - When no process is inside a critical section, any process requesting permission to enter must be allowed to do so immediately
- No process requiring access to its critical section can be delayed forever
- No assumptions are made about relative the speed of processes

# Critical Sections and Mutual Exclusion



# Disabling Interrupts

```
void Extract(int acc_no, int sum)
{
    CLI ();
    int B = Acc[acc_no];
    Acc[acc_no] = B - sum;
    STI ();
}
```

Works only on single-processor systems, but not with user level threads.

Misbehaving/buggy processes may never release CPU

- Mechanism usually only used by kernel code

# Software Solution – Strict Alternation

**P<sub>0</sub>**

turn **0**

**P<sub>1</sub>**

```
while (true) {  
    while (turn != 0)  
        /* loop */ ;  
    critical_section()  
    turn = 1;  
    noncritical_section0();  
}
```

```
while (true) {  
    while (turn != 1)  
        /* loop */ ;  
    critical_section()  
    turn = 0;  
    noncritical_section1();  
}
```

What happens if P<sub>0</sub> takes a long time in its non-critical section?

- Remember: No process running outside its critical section may prevent other processes from entering the critical section

Can we have P<sub>1</sub> execute its loop twice in a row (w/o P<sub>0</sub> executing in-between)?

# Busy Waiting


Strict alternation solution requires continuously testing the value of a variable

Called **busy waiting**

- Wastes CPU time
- Should only be used when the wait is expected to be short

# Atomic Operations

```
void Extract(int acc_no,  
             int sum)  
{  
    int B = Acc[acc_no];  
    Acc[acc_no] = B - sum;  
}
```



```
void Extract(int acc_no,  
             int sum)  
{  
    Acc[acc_no] -= sum;  
}
```

Does this work?

- Not atomic!

Atomic operation: a sequence of one or more statements that is/appears to be indivisible

# Lock Variables

L= 0    lock open/free  
L=1    locked

```
void Extract(int acc_no, int sum)
{
    lock(L) ;
    int B = Acc[acc_no] ;
    Acc[acc_no] = B - sum;
    unlock(L) ;
}
```

```
void lock(int L)
{
    while (L != 0)
        /* wait */ ;
    L = 1;
}
```

```
void unlock(int L)
{
    L = 0;
}
```

- Does this work?

# TSL (Test and Set Lock) Instruction

**Atomic** instruction provided by most CPUs

## TSL (LOCK)

- Atomically sets memory location **LOCK** to 1 and returns old value

### Pseudocode

```
void lock(int L)
{
    while (TSL(L) != 0)
        /* wait */ ;
}
```

### Assembler

TSL L	Read L and set condition code if L=0
BNZ	jumps if Z is not set.
MOV #n,L	sets L to constant n

```
LOCK:    TSL L
         BNZ LOCK
```

```
UNLOCK:  MOV #0, L
```



# Spin Locks

Locks using busy waiting are called **spin locks**

Waste CPU

- Should only be used when the wait is expected to be short

May run into *priority inversion problem*

# Priority Inversion Problem and Spin Locks

## Two processes:

- H with high priority
- L with low priority
- H should always be scheduled if runnable

## Assume the following scenario:

- H is waiting for I/O
- L acquires lock A and enters critical section
- I/O arrives and H is scheduled
- H tries to acquire lock A that L is holding

## What happens?

# Lock Granularity

```
void Extract(int acc_no, int sum)
{
    lock(L) ;
    int B = Acc[acc_no] ;
    Acc[acc_no] = B - sum ;
    unlock(L) ;
}
```

T1: Extract(1, 40) ;

T2: Extract(2, 40) ;

What happens if there are concurrent accesses to *different* accounts?

# Lock Granularity

```
void Extract(int acc_no, int sum)
{
    lock(L[acc_no]);
    int B = Acc[acc_no];
    Acc[acc_no] = B - sum;
    unlock(L[acc_no]);
}
```

T1: Extract(1, 40);

T2: Extract(2, 40);

**Lock granularity:** the amount of data a lock is protecting

*Is finer granularity always better?*

# Lock Overhead and Lock Contention

**Lock overhead:** a measure of the cost associated with using locks

- Memory space
- Initialization
- ***Time required to acquire and release locks***

**Lock contention:** a measure of the number of processes waiting for a lock

- More contention, less parallelism

- **Coarser granularity:**

- Lower overhead
- More contention
- Lower complexity

- **Finer granularity:**

- Higher lock overhead
- Less contention
- Higher complexity

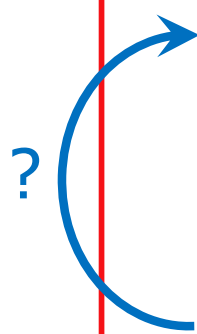
# Minimizing Lock Contention/Maximizing Concurrency

Choose finer lock granularity

- But understand tradeoffs

Release a lock as soon as it is not needed

- Make critical sections small!



```
void AddAccount(int acc_no, int balance)
{
    lock(L_Acc);
    CreateAccount(acc_no);
    lock(L[acc_no]);
    Acc[acc_no] = balance;
    unlock(L[acc_no]);
    unlock(L_Acc);
}
```

# Read/Write Locks

```
void ViewHistory(int acc_no)
{
    print_transactions(acc_no);
}
```

**P1:** ViewHistory(1234) ;

**P2:** ViewHistory(1234) ;

**P3:** ViewHistory(1234) ;

Any locks needed?

# Race Condition

Occurs when multiple threads or processes read and write **shared data** and the final result depends on the relative timing of their execution

- i.e. on the exact process or thread **interleaving**

E.g., the **Extract** example → final value of account 8,000 or 9,000



# Thread Interleavings

```
int a, b; // shared
void P1()
{
    a = 1;
    b = 1;
}
void P2()
{
    b = 2;
    a = 2;
}
```

a = 1	a = 1	a = 1	b = 2	b = 2	b = 2
b = 1	b = 2	b = 2	a = 2	a = 1	a = 1
b = 2	b = 1	a = 2	a = 1	a = 2	b = 1
a = 2	a = 2	b = 1	b = 1	b = 1	a = 2
(2, 2)	(2, 1)	(2, 1)	(1, 1)	(2, 1)	(2, 1)

# Thread Interleaving

Menti.com Q1-3 76 84 50

Consider the following three threads:

T1: {a=1; b=2;}    T2: {b =1;}    T3: {a=2;}

1. How many different thread interleavings are there?
2. If all thread interleavings are as likely to occur, what is the probability to have a=1 and b=1 after all threads complete execution?
3. What about a=2 and b=2?

# Semaphores

Blocking synchronization mechanism invented by Dijkstra in 1965

Idea: Processes will cooperate by means of *signals*

- A process will block, waiting for a specific signal
- A process will continue if it has received a specific signal

**Semaphores** are *special variables*, accessible via the following *atomic* operations:

- **down(s)** : receive a signal via semaphore **s**
- **up(s)** : transmit a signal via semaphore **s**
- **init(s, i)** : initialise semaphore **s** with value **i**


**down()** also called **P()** (*probeer te verlagen*)

**up()** also called **V()** (*verhogen*)

# Semaphores

Semaphores have two private components:

- A counter (non-negative integer)
  - A queue of processes currently waiting for that semaphore
- Queue is typically first in first out (FIFO)



Semaphore  
Data Structure

Queue of processes waiting on Semaphore

# Semaphore Operations

```
init(s, i) ::= counter(s) = i  
              queue(s) = {}
```

```
down(s) ::= if counter(s) > 0  
             counter(s) = counter(s) - 1  
           else  
             add P to queue(s)  
             suspend current process P
```

```
up(s) ::= if queue(s) not empty  
           resume one process in queue(s)  
         else  
           counter(s) = counter(s) + 1
```

# Semaphores for Mutual Exclusion

**Binary semaphore:** counter is initialized to 1

Similar to a lock/mutex

```
process A                                process B
...                                     ...
down(s)                                down(s)
    critical section                    critical section
up(s)                                  up(s)
end                                    end

main() {
    var s:Semaphore
    ...
    init(s, 1) /* initialise semaphore */
    ...
    start processes A and B in random order
    ...
}
```

Note: for binary semaphore if  $s = 1$ ,  $up(s)$  leaves  $s = 1$

# General Semaphores

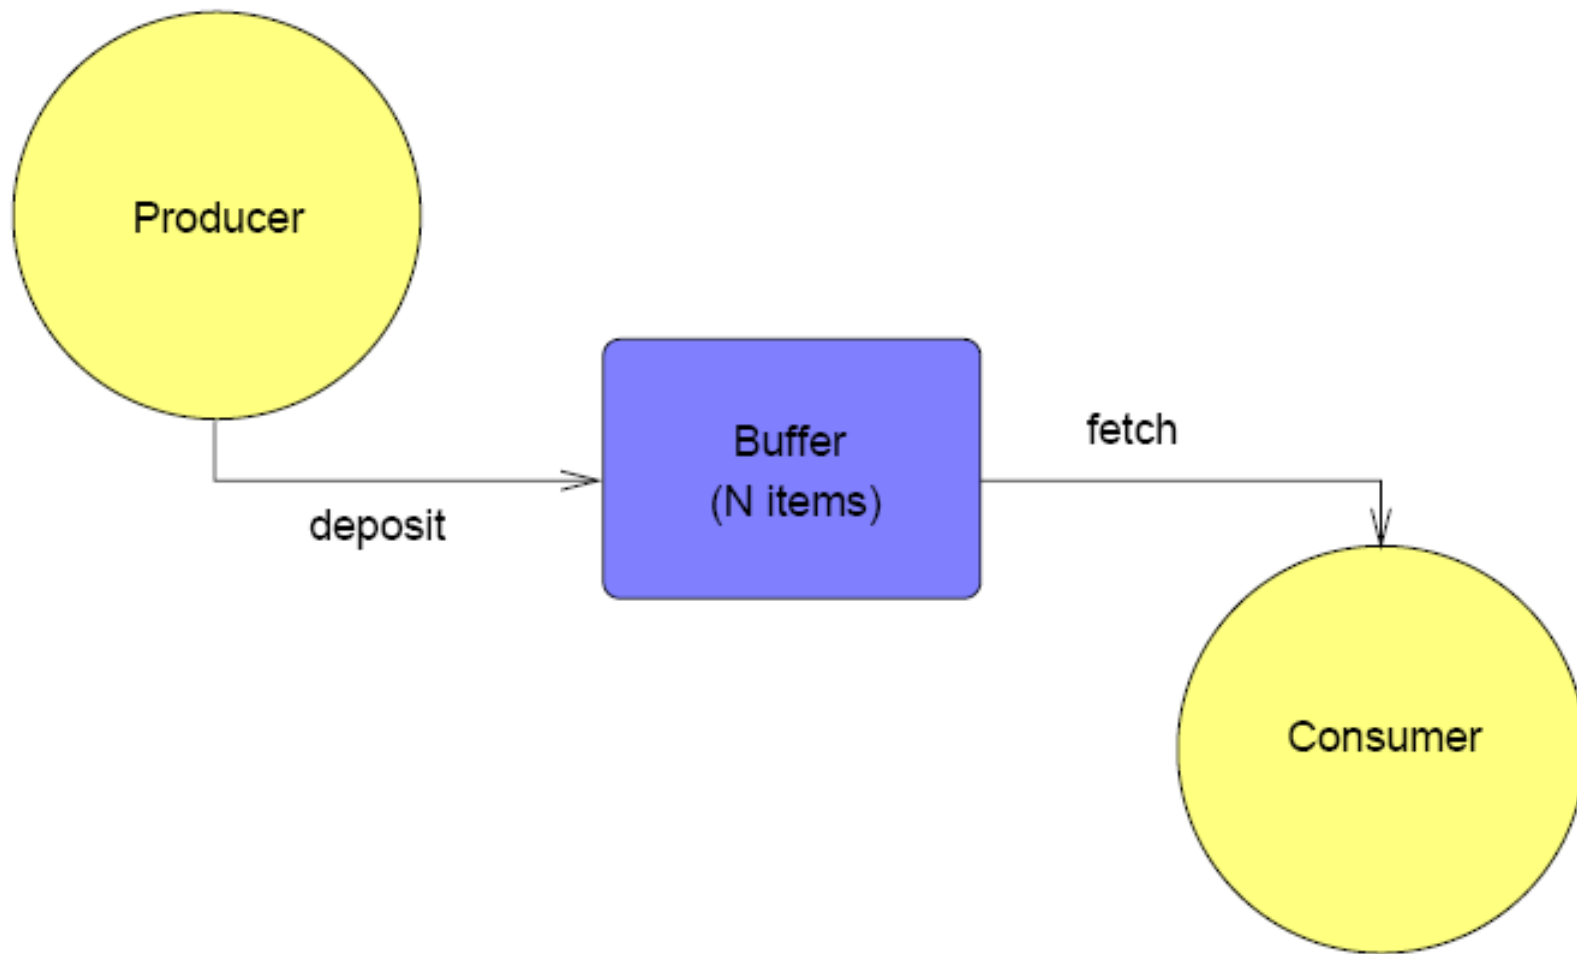
The initial value of a semaphore counter indicates how many processes can access shared data at the same time

**counter(s) >= 0:**

Initial value defines how many processes can execute down without being blocked

Menti.com Q4, 5 threads & semaphore 76 84 50

# Producer / Consumer



There can be multiple producers and consumers



# Producer / Consumer

## Buffer constraints:

- Buffer can hold between **0** and **N** items

## Producer constraints:

- Items can only be deposited in buffer if there is space (items in buffer  $< N$ )
- Items can only be deposited in buffer if mutual exclusion is ensured

## Consumer constraints:

- Items can only be fetched from buffer if it is not empty (items in buffer  $> 0$ )
- Items can only be fetched from buffer if mutual exclusion is ensured

# Producer/Consumer?

```
var item, space, mutex: semaphore
init (item, 0)      /* Semaphore to ensure buffer is not empty */
init (space, N)     /* Semaphore to ensure buffer is not full */
init (mutex, 1)     /* Semaphore to ensure mutual exclusion */
```

```
process Producer
loop
  produce item
  down(mutex)
  down(space)
  deposit item
  up(item)
  up(mutex)
end loop
end Producer
```

```
process Consumer
loop
  down(mutex)
  down(item)
  fetch item
  up(space)
  up(mutex)
  consume item
end loop
end Producer
```

What is wrong with this?

# Producer/Consumer

```
var item, space, mutex: semaphore
init (item, 0)      /* Semaphore to ensure buffer is not empty */
init (space, N)     /* Semaphore to ensure buffer is not full */
init (mutex, 1)     /* Semaphore to ensure mutual exclusion */
```

```
process Producer
```

```
  loop
```

```
    produce item
```

```
    down (space)
```

```
    down (mutex)
```

```
    deposit item
```

```
    up (mutex)
```

```
    up (item)
```

```
  end loop
```

```
end Producer
```

```
process Consumer
```

```
  loop
```

```
    down (item)
```

```
    down (mutex)
```

```
    fetch item
```

```
    up (mutex)
```

```
    up (space)
```

```
    consume item
```

```
  end loop
```


```
end Producer
```

Works for multiple producers & consumers

What happens when space = 0 or items = 0?

Animation: <https://www.youtube.com/watch?v=NuvAjMk9bZ8>

# Readers/Writers



- Writers constraints:
  - items can only be written if no other process is writing;
  - items can only be written if no other process is reading.
- 
- Readers constraints:
  - items can only be read if no other process is writing;
  - items can be read if there are other processes reading.
- 
- File can hold an arbitrary number of items.

# Readers/Writers With Semaphores

```
semaphore mutex, wrt;  
int read_cnt = 0;  
init(mutex, 1);  
init(wrt, 1);
```

```
process writer()  
    loop  
        produce item  
        down(wrt);  
        write item  
        up(wrt);  
    end loop  
end writer
```

Does this work?

```
process reader()  
    loop  
        if(read_cnt == 0)  
            //1st reader  
            down(wrt);  
        down(mutex)  
        read_cnt += 1;  
        up(mutex);  
        read item  
        down(mutex);  
        read_cnt -= 1  
        up(mutex);  
        If (read_cnt == 0)  
            up(wrt);  
        consume item  
    end loop  
end reader
```

# Readers/Writers With Semaphores

```
semaphore mutex, wrt;  
int read_cnt = 0;  
init(mutex, 1);  
init(wrt, 1);
```

```
process writer()  
    loop  
        produce item  
        down(wrt);  
        write item  
        up(wrt);  
    end loop  
end writer
```

Is this fair?

```
process reader()  
    loop  
        down(mutex)  
        read_cnt += 1;  
        if(read_cnt == 1)  
            //1st reader  
            down(wrt);  
        up(mutex);  
        read item  
        down(mutex);  
        read_cnt -= 1  
        If (read_cnt == 0)  
            up(wrt);  
        up(mutex);  
        consume item  
    end loop  
end reader
```

# Semaphore Question

The following program consists of 3 concurrent processes and 3 binary semaphores. The semaphores are initialized as  $S0 = 1$ ,  $S1 = 0$ ,  $S2 = 0$ .

Process P0

```
while true
{ down(S0) ;
  print '0' ;
  up(S1) ;
  up(S2) }
```

Process P1

```
down(S1) ;
up(S0) ;
```

Process P2

```
down(S2) ;
up(S0) ;
```

How many times will P0 print '0'?

Menti.com Q6 76 84 50

# General Semaphore Using Binary Semaphores

Describe a suitable data structure for a general semaphore and give a pseudocode outline for the following operations in terms of the operations `down(s)` and `up(s)` on a binary semaphore `s`.

<code>init (value, gs)</code>	initialises the general semaphore <code>gs</code> to <code>value</code>
<code>gen_down (var gs)</code>	the down operation on a general semaphore <code>gs</code>
<code>gen_up (var gs)</code>	the up operation on a general semaphore <code>gs</code>



# Monitors

Higher-level synchronization primitive

Introduced by Hansen (1973) and Hoare (1974)

Refined by Lampson (1980)

# Monitors

Ensure mutual exclusion for shared resource (data)

Entry procedures

- Can be called from outside the monitor

Internal procedures

- Can be called only from monitor procedures

An (implicit) monitor lock

One or more condition variables

Processes can only call entry procedures

- cannot directly access internal data

Only one process can be in the monitor at one time

# Condition Variables

## Associated with high-level conditions

- “some space has become available in the buffer”
- “some data has arrived in the buffer”

## Operations:

- **wait(c)**: releases monitor lock and waits for **c** to be signalled
- **signal(c)**: wakes up one process waiting for **c**
- **broadcast(c)**: wakes up all processes waiting for **c**

## Signals do not accumulate i.e c is not a counter.

- If a condition variable is signalled with no one waiting for it, the signal is lost

# What happens on signal?

[Hoare] A process waiting for signal is immediately scheduled

- + Easy to reason about
- Inefficient: the process that signals is switched out, even if it has not finished yet with the monitor
- Places extra constraints on the scheduler

[Lampson] Sending signal and waking up from a wait are not atomic

- More difficult to understand, need to take extra care when waking up from a wait()
- + More efficient, no constraints on the scheduler
- + More tolerant of errors: if the condition being notified is wrong, it is simply discarded when rechecked (see next slides)

*Usually [**Lampson**] is used*

# Hoare Monitor Implementation Using Semaphores

## Variables

```
semaphore mutex;    // (initially = 1)
semaphore next;     // (initially = 0)
int next_count = 0;
```

Each access procedure ***F*** will be replaced by

```
wait(mutex) ;
...
body of access procedure F;
...
if (next_count > 0)
    up(next)
else
    up(mutex) ;
```

Mutual exclusion within a monitor is ensured

Note: this code is generated by a compiler and is not seen by the programmer who writes the access procedures.

# Monitor Implementation – Condition Variables

For each condition variable **c**, we have:

```
semaphore c_sem; // (initially = 0)
int c_count = 0;
```

The operation **wait (c)** can be implemented as:

```
c_count++;
if (next_count > 0) up(next);
else up(mutex);
down(c_sem);
c_count--;
```

The operation **signal (c)** can be implemented as:

```
if (c_count > 0) {
    next_count++; up(c_sem);
    down(next); next_count--;
}
```

# Producer/Consumer with Monitors

```
monitor ProducerConsumer
    condition not_full, not_empty;
    integer count = 0;

    entry procedure insert(item)
        if (count == N) wait(not_full);
        insert_item(item); count++;
        signal(not_empty);

    entry procedure remove(item)
        if (count == 0) wait(not_empty);
        remove_item(item); count--;
        signal(not_full);
end monitor
```

**Does this work?**

# Producer/Consumer with Lampson Monitors

```
monitor ProducerConsumer
    condition not_full, not_empty;
    integer count = 0;

    entry procedure insert(item)
        while (count == N) wait(not_full);
        insert_item(item); count++;
        signal(not_empty);

    entry procedure remove(item)
        while (count == 0) wait(not_empty);
        remove_item(item); count--;
        signal(not_full);
end monitor
```



# Readers/Writers Revisited

## Correctness Constraints:

- Readers can access file when no writers
- Writers can access file when no readers or writers
- Only one thread manipulates state variables at a time

## Basic structure of a solution:

- Reader()
  - Wait until no writers
  - Access file
  - Check out – wake up a waiting writer
- Writer()
  - Wait until no active readers or writers
  - Access file
  - Check out – wake up waiting readers or writer

# Readers/Writers: Fairness?

## Problem statement clarification

- Suppose that a writer is active and a mixture of readers and writers now shows up. Who should get in next?
- If a writer is waiting and an endless stream of readers keeps showing up. Is it fair for them to become active?

## Alternation is a possible fair solution:

- Once a reader is waiting, readers will get in next.
- If a writer is waiting, one writer will get in next.

## State variables needed (Protected by a lock called "lock"):

- int NReaders: Number of active readers; initially = 0
- int WaitReaders: Number of waiting readers; initially = 0
- int N Writers: Number of active writers; initially = 0
- int WaitWriters: Number of waiting writers; initially = 0
- Condition CanRead = NIL, CanWrite = NIL

# Readers/Writers with Monitors

```
monitor ReadersNriters
  integer WaitWriters, WaitReaders,
           NReaders, N Writers;
  condition CanRead, CanWrite;

  entry procedure StartRead()
    if(NWriters == 1 or WaitWriters > 0)
    {
      ++WaitReaders; Wait(CanRead); --WaitReaders;
    }
    ++Nreaders;
    Signal(CanRead);
  end StartRead

  entry procedure EndRead()
    If(--Nreaders == 0) Signal(CanWrite);
  end EndRead
```

# Reader/Writer contd

```
entry procedure StartWrite()  
    if(NWriters == 1 or NReaders > 0)  
    {  
        ++WaitWriters; wait(CanWrite); --WaitWriters;  
    }  
    NWriters = 1;  
end StartWrite;
```

```
entry procedure EndWrite()  
    NWriters = 0;  
    if(WaitReaders > 0) Signal(CanRead);  
    else Signal(CanWrite);  
end EndWrite;
```

```
end monitor
```

# Monitors

Monitors are a language construct

Not supported by C

Java

- synchronized methods
- no condition variables
  - ✦ wait() and notify()

# Synchronization within Monitors

Synchronization within monitors uses condition variables and two special operations, **wait** and **signal**. A more general form of synchronization would be to have a single primitive, **waituntil**, that had an arbitrary Boolean predicate as parameter. Thus, one could say, for example,

**waituntil**      $(x < 0 \text{ or } y + z < n)$

The signal primitive would be no longer needed. This scheme is clearly more general than that of Hoare, but it is not used. Why not?

(Hint: think about the implementation.)

# Bohr and Heisen bugs

## Bohrbugs:

- Deterministic, reproducible bugs
- Behave similar to Bohr's atom model where electrons deterministically orbit the nucleus

## Heisenbugs

- Non-deterministic, hard to reproduce bugs
  - ✦ Often caused by race conditions
- Suffer from the observer effect (Heisenberg Uncertainty Principle): attempts to observe them (i.e., printf's) make them disappear!

## Which bug would you rather have?

- During development/testing: \_\_\_\_\_
- During deployment: \_\_\_\_\_

# Communication & Synchronization Summary

Signals: really interaction with kernel, to wake a waiting process or indicate a problem.

Pipes: simple read, write type communication

Shared memory: requires synchronisation to prevent corruption

Critical section: code in which process accesses shared resource

Mutual exclusion: only 1 process at a time within CS

Disabling interrupts: may not be effective

Locks: low level, busy wait, very difficult to program correctly

Semaphores: blocks waiting program, but difficult to program

Monitors: easier to program, but signal semantics can be tricky