

# Memory Management

Anandha Gopalan

(with thanks to R. Kolcun and P. Pietzuch)

[axgopala@imperial.ac.uk](mailto:axgopala@imperial.ac.uk)

### Basic Concepts

- Memory Allocation
- Swapping

### Virtual Memory

### Paging & Segmentation

- Demand Paging
- Page replacement algorithms
- Working set model


### Linux Memory Management

# Memory Hierarchy

Hardware: CPU registers and main memory

- Register access in one CPU clock cycle (or less)
- Main memory can take many cycles
- Caches sit between main memory and CPU registers

Managed by




# Memory Hierarchy

Hardware: CPU registers and main memory

- Register access in one CPU clock cycle (or less)
- Main memory can take many cycles
- Caches sit between main memory and CPU registers

Managed by



Memory is a key component of the computer

- e.g. every instruction cycle involves memory access  $\Rightarrow$  process has to be loaded into memory before it can execute

Memory management needs to provide

- Memory allocation
- Memory protection

Characteristics

- No knowledge of how memory addresses are generated
  - e.g. instruction counter, indexing, indirection, ...
- No knowledge what memory addresses are used for
  - e.g. instructions or data
- True for simple case but may want protection with respect to read, write, execute, etc.

# Logical vs. Physical Address Space

Memory management binds logical address space to physical address space

## Logical address

- Generated by the CPU
- Address space seen by the process

## Physical address

- Address seen by the memory unit
- Refers to physical system memory

Logical and physical addresses

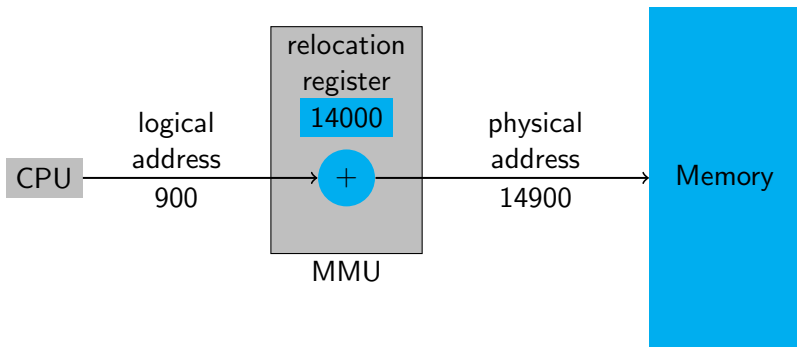
- Same in compile- and load-time address-binding schemes
- Different in execution-time address-binding schemes

How do you achieve this mapping?

# Memory-Management Unit (MMU)

Hardware device for mapping logical to physical addresses

- e.g. add value in relocation register to every address generated by process when sent to memory
- User process deals with logical addresses only
- Has to be fast → implemented in hardware



Main memory is usually split into two partitions:

- Resident operating system (**kernel**)
  - Usually held in low memory with interrupt vector
- User processes (**user**)
  - Held in high memory

How do you decide where to load a new process?

Need to figure out the strategy for process to be loaded into the correct location




## Contiguous allocation with **relocation** registers

- **base** register contains physical start address for process
- **limit** register contains maximum logical address for process
- MMU maps logical address dynamically
  - Physical address = logical address + **base**
  - If logical address > **limit** then error

# Contiguous Memory Allocation II

**base** and **limit** register define logical address space



e.g jmp 100 in program would go to physical location 300140

# Multiple-Partition Allocation

## Hole


- Block of available memory
- Holes of various size scattered throughout memory

When new process arrives:

- allocate memory from hole large enough

OS maintains information about:

- Allocated partitions
- Free partitions (holes)



What is the best algorithm for allocation?

# Dynamic Memory Allocation

**First-fit** → Allocate first hole that is big enough

**Best-fit** → Allocate smallest hole that is big enough

- Must search entire list, unless ordered by size
- Produces smallest leftover hole

**Worst-fit** → Allocate largest hole

- Must also search entire list
- Produces largest leftover hole

## Why best-fit or worst-fit?

First-fit and best-fit better than worst-fit in terms of speed and storage utilisation

# Fragmentation

**External** fragmentation → memory exists to satisfy request, but not contiguous




Reduce external fragmentation by compaction

- Shuffle memory contents to place all free memory together in one large block → leads to I/O bottlenecks

# Fragmentation

**External** fragmentation → memory exists to satisfy request, but not contiguous




Reduce external fragmentation by compaction

- Shuffle memory contents to place all free memory together in one large block → leads to I/O bottlenecks

# Fragmentation

**External** fragmentation → memory exists to satisfy request, but not contiguous




Reduce external fragmentation by compaction

- Shuffle memory contents to place all free memory together in one large block → leads to I/O bottlenecks

# Fragmentation

**External** fragmentation → memory exists to satisfy request, but not contiguous



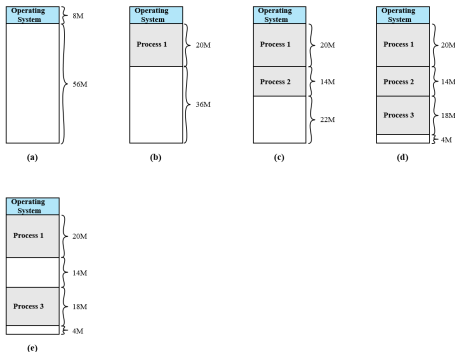
Reduce external fragmentation by compaction

- Shuffle memory contents to place all free memory together in one large block → leads to I/O bottlenecks



# Fragmentation

**External** fragmentation → memory exists to satisfy request, but not contiguous

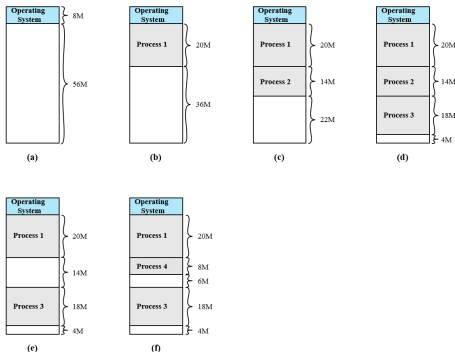


Reduce external fragmentation by compaction

- Shuffle memory contents to place all free memory together in one large block → leads to I/O bottlenecks

# Fragmentation

**External** fragmentation → memory exists to satisfy request, but not contiguous

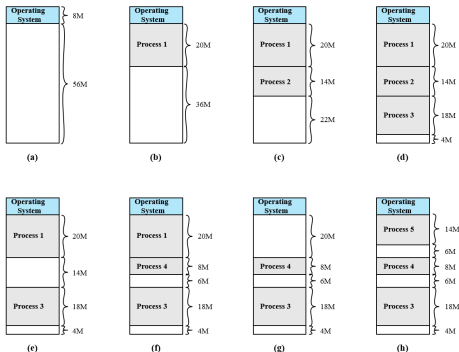


Reduce external fragmentation by compaction

- Shuffle memory contents to place all free memory together in one large block → leads to I/O bottlenecks

# Fragmentation

**External** fragmentation → memory exists to satisfy request, but not contiguous



Reduce external fragmentation by compaction

- Shuffle memory contents to place all free memory together in one large block → leads to I/O bottlenecks


Problem: Number of processes limited by amount of available memory

- But ... only running processes need to be in memory

Solution:

- **Swap** processes temporarily out of memory to backing store
- Bring back into memory for continued execution
- Requires **swap space** → can be file or dedicated partition on disk
- **Transfer time** is major part of swap time

# Swapping




What if a process is “too large” to fit into memory  $\Rightarrow$  can only part of a process exist in memory?

Separation of user logical memory from physical memory

- Only part of process needs to be in memory for execution
- Logical address space can be much larger than physical address space
- Address spaces can be shared by several processes
- Allows for more efficient process creation

# Virtual Memory



Virtual memory can be implemented via

- Paging
- Segmentation



# Paging

Physical address space of process can be noncontiguous

- Process allocated physical memory when available
  - Avoid external fragmentation
  - Avoid problems of variable sized memory chunks

## Frames

- Fixed-sized blocks of physical memory
- Keep track of all free frames

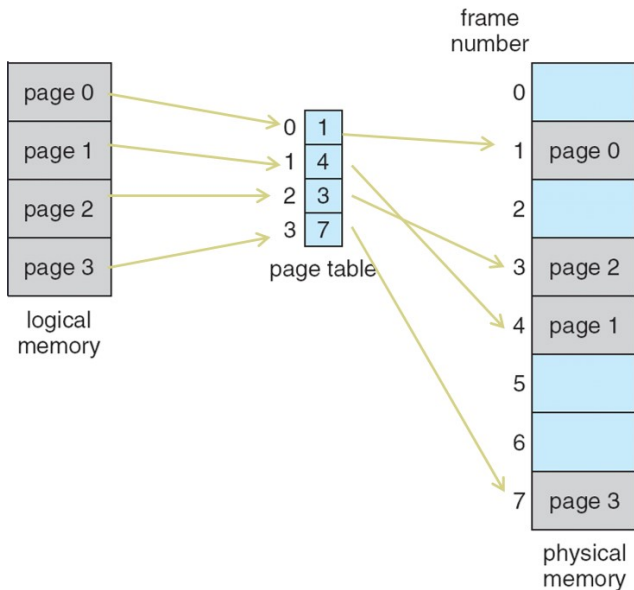
## Pages

- Block of same size (as **frame**) of logical memory

To run program of size  $n$  pages

- Find  $n$  free frames and load program
- Set up page table to translate logical into physical addresses

# Page Table Example



How does logical address translate to physical address?

Hint: pages and frames are the same size  $\Rightarrow$  address offset in the page will be the **same** as that in the frame

Address now consists of two parts: **page number** and **page offset**

- only need to translate page number into its corresponding frame address

## How do you calculate the page number?

Depends on address size and page/frame size

e.g. Consider a page/frame size of 64 bytes

- 64 bytes can be addressed  $\Rightarrow$  total of 64 addresses
- Number of bits required for 64 addresses = 6 ( $2^6 = 64$ )

For a 10-bit **virtual** address we have:

- page offset requires 6 bits (based on above)
- page number has 4 bits (remaining bits)  $\Rightarrow$  between 0 ... 15

# Address Translation III

## Page number (p)

- Used as an index into page table
- Page table has base address of pages in physical memory

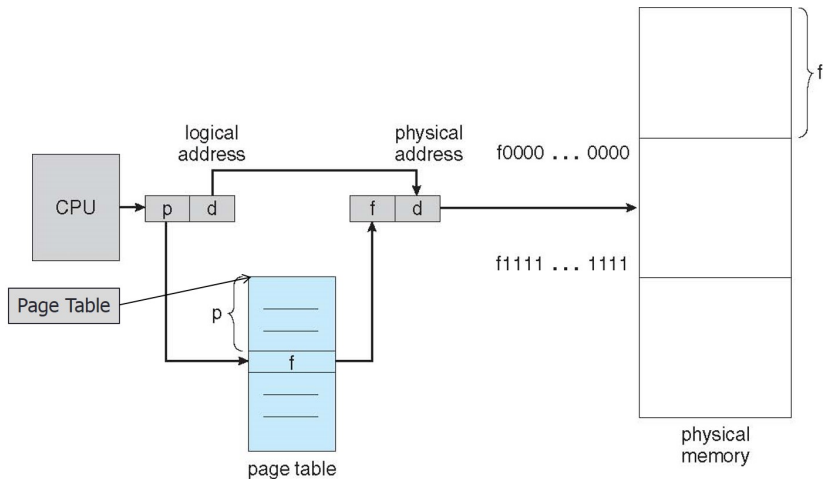
## Page offset (d)

- Defines physical memory address sent to the memory unit
- Combined with base address

For given address size of  $m$ -bits and page size of  $2^n$

page number	page offset
p	d
$(m - n)$ bits	$n$ bits


# Paging Hardware




# Free Frames

free-frame list

14  
13  
18  
20  
15



13  
14  
15  
16  
17  
18  
19  
20  
21




(a)


Before allocation

free-frame list

15



13  
14  
15  
16  
17  
18  
19  
20  
21



(b)

After allocation

0	14
1	13
2	18
3	20

new-process page table

# Example Problem

## Address Translation

Consider a 32-bit virtual memory address and a page size of 1 KB. How many pages can a process potentially have?



## Example Problem

### Address Translation

Consider a 32-bit virtual memory address and a page size of 1 KB.  
How many pages can a process potentially have?

1 KB page size = 1024 bytes  $\Rightarrow$  total of 1024 addresses

Number of bits needed for 1024 address = 10 ( $2^{10} = 1024$ )

# Example Problem

## Address Translation

Consider a 32-bit virtual memory address and a page size of 1 KB.  
How many pages can a process potentially have?

1 KB page size = 1024 bytes  $\Rightarrow$  total of 1024 addresses

Number of bits needed for 1024 address = 10 ( $2^{10} = 1024$ )

For a 32-bit address you have:

- page offset requires 10 bits
- page number has 22 bits  $\Rightarrow 2^{22}$  (4194304) potential pages

# Fragmentation

**Internal** fragmentation  $\rightarrow$  Allocated memory is larger than requested memory, but size difference internal to partition

## Example - Calculating Internal Fragmentation

Page size = 2048 bytes; Process size = 72,766 bytes

Number of pages =  $\frac{72766}{2048} = 35$

Bytes left-over =  $72766 \% 2048 = 1086$

Internal fragmentation =  $2048 - 1086 = 962$  bytes

Worst-case fragmentation  $\Rightarrow$  1 frame = 1 byte

Average-case fragmentation  $\Rightarrow \frac{1}{2}$  frame size

### Are small frames desirable?

- Each page table entry takes memory to track
- Page size growing over time → typically 4 KB but some architectures support variable page sizes up to 256 MB

# Page Table Implementation

Page table kept in main memory

- **Page-table base register** (PTBR) points to page table
- Context switch requires update of PTBR for new process page table (if necessary)
- **Page-table length register** (PTLR) indicates size

Problem

- Inefficient, as every data/instruction access requires two memory accesses → one for page table and one for data/instruction

# Associative Memory

Solution: use special fast-lookup hardware cache as associative memory

Associative memory  $\rightarrow$  supports parallel search

- Called **Translation Look-aside Buffer (TLB)**

Page #	Frame #
1	4
3	74
9	50
7	7

- Address translation (p, d)
  - If p in associative register, get frame # out
  - Otherwise, get frame # from page table in memory

# Translation Look-aside Buffers (TLBs) I


TLBs usually needs to be flushed after context switch

- Can lead to substantial overhead
- What about kernel pages for system calls?

Some TLBs store address-space ids (ASIDs) in entries


- Uniquely identifies each process to provide address-space protection for that process

# Translation Look-aside Buffers (TLBs) II






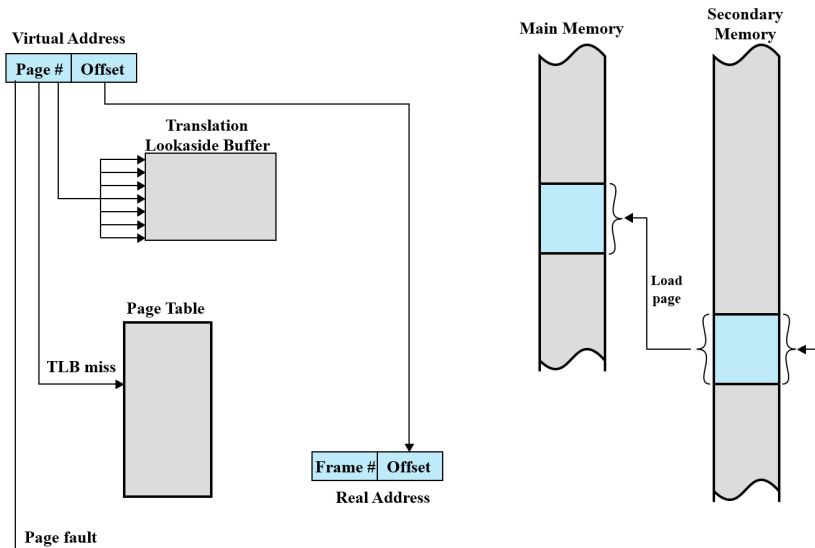
# Translation Look-aside Buffers (TLBs) II




# Translation Look-aside Buffers (TLBs) II



# Translation Look-aside Buffers (TLBs) II



# Translation Look-aside Buffers (TLBs) II



# Example Problem

## Effective Access Time

TLB Lookup =  $\epsilon$  (can be  $< 10\%$  of memory access time  $m$ )

Hit Ratio =  $\alpha$

- Fraction of times that page is found in associative registers
- Ratio related to number of associative registers

$$\begin{aligned}\text{Effective Access Time (EAT)} &= (\epsilon + m) \times \alpha + (\epsilon + 2m) \times (1 - \alpha) \\ &= 2m + \epsilon - m\alpha\end{aligned}$$

Consider  $\alpha = 80\%$ ,  $\epsilon = 10 \text{ ns}$  for TLB search,  $m = 100 \text{ ns}$  for memory access

- $\text{EAT} = 110 \times 0.80 + 210 \times 0.20 = 130 \text{ ns}$

A more realistic hit ratio might be 99%

- $\text{EAT} = 110 \times 0.99 + 210 \times 0.01 = 111 \text{ ns}$

## Why do we need need to worry?

Page table can grow to be very large in size

On a 32-bit machine with a 4 KB page size:

- Number of page table entries =  $\frac{2^{32}}{2^{12}} = 2^{20}$
- Size of each page table entry = 32 bits
- Size of page table =  $2^{20} \times 32 \text{ bits} = 4 \text{ MB}$

On 64-bit machine with 4 KB pages → page table needs  $2^{52}$  entries

- with 8 bytes per entry, that's 30 million GB ...
- lot of memory to be allocated ☹

**Hierarchical** page table

**Hashed** page table

**Inverted** page table

# Hierarchical Page Table I

**Idea:** Let the page-table be broken-up and paged if it is too large

Simple technique → **two-level page table** for a machine with 32-bit addresses and a 4 KB page size

- page offset needs 12 bits
- page table size = 4 MB


How do you break the page table up?

Each part of the page table that is being paged must fit on a page

- Recall: page size = 4 KB
- Number of entries on one page =  $\frac{\text{Page size}}{\text{Address size}} = \frac{4 \text{ KB}}{32 \text{ bits}} = 2^{10}$
- No of bits required for  $2^{10}$  entries = 10
- Address bits left for top-level page table =  $32 - 10 - 12 = 10$



# Hierarchical Page Table II



**Fix** outer page table in memory

# Two-Level Paging I

Logical address divided

- Page number consisting of 20 bits
- Page offset consisting of 12 bits

Since page table paged, page number further divided

- 10 bit page number
- 10 bit page offset within  $2^{nd}$  level page table


Thus, logical addresses as follows

page number		page offset
$p_1$	$p_2$	d
10	10	12

$p_1 \rightarrow$  index into the outer page table

$p_2 \rightarrow$  displacement within page pointed to by outer page table

# Two-Level Paging II



# Example Problem

## Page Table Addressing

Consider a paging system that uses a three-level page table. Virtual addresses are composed into four fields (a, b, c, d) with d being the offset. What is the maximum number of pages in a virtual address space?

Answer:  $2^{a+b+c}$ , since there are a total of  $2^{a+b+c+d}$  addresses in the address space and each page has  $2^d$  addresses

# Example Problem

## Page Table Addressing

Consider a paging system that uses a three-level page table. Virtual addresses are composed into four fields (a, b, c, d) with d being the offset. What is the maximum number of pages in a virtual address space?

Answer:  $2^{a+b+c}$ , since there are a total of  $2^{a+b+c+d}$  addresses in the address space and each page has  $2^d$  addresses

## Page Table: Another Idea


Don't store entry per **page** but per **frame**

- **Hashed page table**
- **Inverted page table**

# Hashed Page Table

Hash virtual page number into page table

- Page table contains chain of elements hashing to same location
- Search for match of virtual page number in chain
- Extract corresponding physical frame if match found

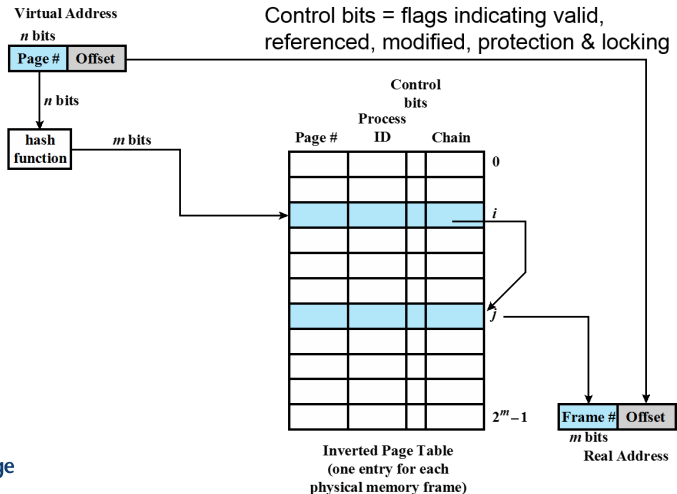


# Inverted Page Table

One entry per physical frame

Decreases memory needed to store page table

- But increases time to search table when page reference occurs



Paging gives one-dimensional virtual address space → what about separate address spaces for code, data, stack?

## Segment


- Independent address space from 0 to some maximum
- Can grow/shrink independently
- Support different kinds of protection (read/write/execute)
- Unlike pages, programmers are aware of segments
- Segment corresponds to program, procedure, stack, object, array, etc.

Memory allocation harder due to variable size


- May need to move segment which grows
- May suffer from external fragmentation
- But good for shared libraries



# Logical View of Segmentation

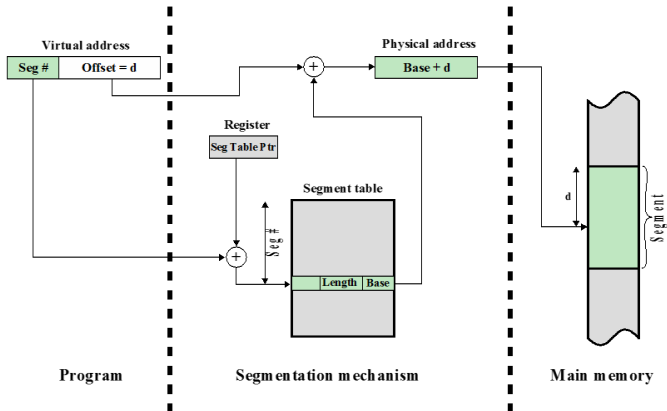


User logical space



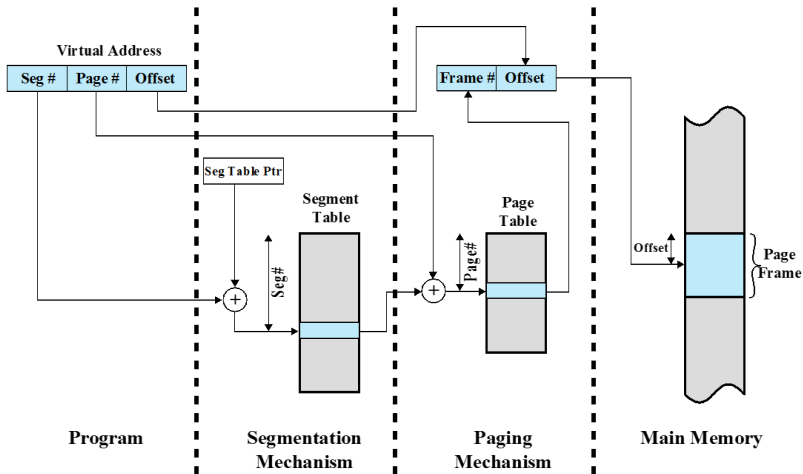
Physical memory space

# Segmentation Address Translation



- One bit in table indicates whether segment is in memory
- Another bit indicates whether segment is modified

# Hybrid Segmentation/Paging



Most OSs use only paging




**Protection bits** → associated with a frame indicate read-only, read-write, execute only

## **Valid-invalid bit**

- **Valid** → page present in physical memory
- **Invalid** → page missing in physical memory
  - **Page fault** is generated  $\Rightarrow$  kernel trap to bring in page from backing store

**Page replacement bits** → to indicate if page has been modified or referenced (used later). Also, lock bit to prevent page from being transferred out

# Memory Validity



When do you bring the page into memory?

Bring page into memory only **when needed**

- Lower I/O load
- Less memory needed
- Faster response time
- Support for more users


Page needed → reference it

- Invalid reference → abort
- Not-in-memory → bring into memory

Many page faults when process first starts

Eventually required pages are in memory so page fault rate drops

# Demand Paging II




Use **valid-invalid** bit to check memory validity

- 1 → in memory
- 0 → not in memory
  - Initially set to 0 on all entries
  - If 0 during address translation → page fault



# Demand Paging IV



First reference, trap to OS → page fault


OS looks at another table to decide

- Invalid reference → abort
- Valid reference but just not in memory → handle request

To handle valid request

- Get empty frame
- Swap page into frame
- Reset tables, validation bit = 1
- Restart last instruction

# Page Faults II



# Performance: Demand Paging

Page Fault Rate ( $p$ ),  $0 \leq p \leq 1.0$

- If  $p = 0$ , no page faults
- If  $p = 1$ , every reference causes a page fault

$$\begin{aligned} \text{Effective Access Time (EAT)} &= (1 - p) \times \text{memory access} + p \\ &\times (\text{page fault overhead} \\ &+ [\text{swap page out}] \\ &+ \text{swap page in} \\ &+ \text{restart overhead}) \end{aligned}$$

Note: no need to swap page out if not modified

## Copy-on-Write (COW)

- Allows parent and child processes to initially share same pages in memory → if either process modifies shared page, then copy page
- Efficient process creation: copy only modified pages
- Free pages allocated from pool of zeroed-out pages

## Memory-mapped files

- Map file into virtual address space using paging
- Simplifies programming model for I/O

## I/O Interlock

- Pages must sometimes be locked into memory
  - Pages used for DMA from disk

# Example Problem

## Demand Paging

Memory access time = 200 *ns*

Average page-fault service time = 8 *ms*

$$\begin{aligned}\text{EAT} &= (1 - p) \times 200 + p \times (8 \text{ ms}) \\ &= (1 - p) \times 200 + p \times 8,000,000 \\ &= 200 + p \times 7,999,800\end{aligned}$$

If one access out of 1,000 causes a page fault, then  $\text{EAT} = 8.2 \text{ ms} \rightarrow$  slowdown by a factor of 40!

If we want performance degradation < 10%

$$\begin{aligned}\text{EAT} &< \text{EAT} + 10\% \text{ of EAT} \\ 200 + 7,999,800 \times p &< 220 \\ 7,999,800 \times p &< 20 \\ p &< 0.0000025\end{aligned}$$

Less than one page fault in every 400,000 memory accesses

# Page Replacement

No free frame? Replace page

How do you decide which page to replace?

Find some unused page in memory to swap out → need strategy for **page replacement**

Minimise number of page faults

- Avoid bringing same page into memory several times

Prevent over-allocation of memory

- Page-fault service routine should include page replacement

Use modify (dirty) bit to reduce overhead of page transfers



- Only modified pages written to disk

# Basic Page Replacement I

- 1 Find location of desired page on disk
- 2 Find free frame
- 3 Frame found?
  - Yes?  $\Rightarrow$  use it
  - No?  $\Rightarrow$  use replacement algorithm to select **victim** frame
- 4 Load desired page into (newly) freed frame
- 5 Update page and frame tables
- 6 Restart process



# Basic Page Replacement II



# Page Replacement Algorithms

**Aim:** Lowest page-fault rate

How do we compare page replacement algorithms?

Use a **Reference String** → particular string of memory references and calculate number of page faults for each algorithm

E.g. 1, 2, 3, 3, 2, 4, 1, 4, 5, 5, 7, 2, 3, 1

# Optimal Algorithm

Replace page that will not be used for the longest period of time

- Unimplementable, as knowledge of future references needed
- Useful for measuring how well other algorithms perform



Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Optimal Algorithm

Replace page that will not be used for the longest period of time

- Unimplementable, as knowledge of future references needed
- Useful for measuring how well other algorithms perform

1

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Optimal Algorithm

Replace page that will not be used for the longest period of time

- Unimplementable, as knowledge of future references needed
- Useful for measuring how well other algorithms perform

1
2

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Optimal Algorithm

Replace page that will not be used for the longest period of time

- Unimplementable, as knowledge of future references needed
- Useful for measuring how well other algorithms perform

1
2
3

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Optimal Algorithm

Replace page that will not be used for the longest period of time

- Unimplementable, as knowledge of future references needed
- Useful for measuring how well other algorithms perform

1
2
3
4

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Optimal Algorithm

Replace page that will not be used for the longest period of time

- Unimplementable, as knowledge of future references needed
- Useful for measuring how well other algorithms perform

1
2
3
4

 5

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# Optimal Algorithm

Replace page that will not be used for the longest period of time

- Unimplementable, as knowledge of future references needed
- Useful for measuring how well other algorithms perform

1	4
2	
3	
4	5

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Optimal Algorithm

Replace page that will not be used for the longest period of time

- Unimplementable, as knowledge of future references needed
- Useful for measuring how well other algorithms perform

1	4
2	
3	
4	5

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

6 page faults

# Optimal Algorithm

Replace page that will not be used for the longest period of time

- Unimplementable, as knowledge of future references needed
- Useful for measuring how well other algorithms perform

1	4
2	
3	
4	5

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

6 page faults

Reference string

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2		2		2			2			2				7		
	0	0	0		0		4			0			0				0		
		1	1		3		3			3			1				1		

Total of 9 page faults

# First-In-First-Out (FIFO) Algorithm

Replace oldest page

- May replace heavily used page

Reference string

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2		2	2	4	4	4	0			0	0			7	7	7
	0	0	0		3	3	3	2	2	2			1	1			1	0	0
		1	1		1	0	0	0	3	3			3	2			2	2	1

Total of 15 page faults

Heavily used pages, 0, 2, 3 are being swapped in and out

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)



Assume 3 frames



Assume 4 frames

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)



Assume 3 frames



Assume 4 frames

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1
2

Assume 3 frames

1
2

Assume 4 frames

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1
2
3

Assume 3 frames

1
2
3

Assume 4 frames



# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1
2
3

4

Assume 3 frames

1
2
3
4

Assume 4 frames

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1	4
2	1
3	

Assume 3 frames

1
2
3
4

Assume 4 frames

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1	4
2	1
3	2

Assume 3 frames

1
2
3
4

Assume 4 frames

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1	4	5
2	1	
3	2	

Assume 3 frames

1	5
2	
3	
4	

Assume 4 frames

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1	4	5
2	1	
3	2	

Assume 3 frames

1	5
2	1
3	
4	

Assume 4 frames

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1	4	5
2	1	
3	2	

Assume 3 frames

1	5
2	1
3	2
4	

Assume 4 frames

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1
2
3

4 5  
1 3  
2

Assume 3 frames

1
2
3
4

5  
1  
2  
3

Assume 4 frames

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1	4	5
2	1	3
3	2	4

Assume 3 frames

1	5	4
2	1	
3	2	
4	3	

Assume 4 frames



# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1	4	5	Assume 3 frames 9 page faults
2	1	3	
3	2	4	

1	5	4	Assume 4 frames 10 page faults
2	1	5	
3	2		
4	3		

# Belady's Anomaly

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 (FIFO replacement)

1	4	5	Assume 3 frames 9 page faults
2	1	3	
3	2	4	

1	5	4	Assume 4 frames 10 page faults
2	1	5	
3	2		
4	3		

Belady's Anomaly: More frames  $\Rightarrow$  more page faults

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter



Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1
2

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1
2
3

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1
2
3
4

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1
2
3
4

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1
2
3
4

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1	5
2	
3	
4	

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1	5
2	
3	
4	

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1	5
2	
3	
4	

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1	
2	
3	5
4	3

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1
2
3
4

5 4  
3

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1	5
2	
3	5
4	3

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

8 page faults

# Least Recently Used (LRU) Algorithm

Each page entry has a counter

- When page referenced, copy clock into counter
- When page needs to be replaced, choose lowest counter

1	5
2	
3	5
4	3

Assume 4 frames

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

8 page faults

Reference string

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

Total of 12 page faults



# LRU Approximation Algorithms

Proper LRU is expensive  $\rightarrow$  use approximations instead

## Reference bit

- With each page associate reference bit  $r$ , initially  $r = 0$ 
  - When page referenced, set  $r = 1$
  - Replace page with  $r = 0$  (if one exists)
- Periodically reset reference bits
- Does not provide proper order for LRU



## Clock Replacement Policy

- Needs reference bit  $r$  and uses clock replacement
- If page to be replaced (in clock order) has  $r = 1$  then
  - Set  $r = 0$  and leave page in memory
  - Continue till you find  $r = 0$ , and replace that page
  - If all  $r = 1$ , replace starting page

# Clock Page Replacement

When page fault occurs, the page being pointed to is inspected

- If  $r = 0$ , evict page
- If  $r = 1$ , clear  $r$ , and advance pointer



Keep counter of number of references made to each page

## **LFU (least frequently used) algorithm**

- Replace page with smallest count
- May replace page just brought into memory
- Page with heavy usage in past will have high count
  - Reset counters or use **aging**

## **MFU (most frequently used) algorithm**

- Replace page with largest count
- Page with smallest count probably just brought in and yet to be used

# Example Problem

## Page Replacement

Reference string: 1, 2, 1, 3, 2, 1, 4, 3, 1, 1, 2, 4, 1, 5, 6, 2, 1.

Assuming number of frames is 3, calculate the number of page faults for LRU and Clock page replacement algorithms.

Using LRU:

1	2	1	3	2	1	4	3	1	1	2	4	1	5	6	2	1
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	2	2
	2	2	2	2	2	2	3	3	3	3	4	4	4	4	6	6
			3	3	3	4	4	4	4	2	2	2	5	5	5	1
Y	Y	N	Y	N	N	Y	Y	N	N	Y	Y	N	Y	Y	Y	Y

Total of 11 page faults

Using Clock:

1	2	1	3	2	1	4	3	1	1	2	4	1	5	6	2	1
1	1	1	1	1	1	4	4	4	4	4	4	4	5	5	5	5
	2	2	2	2	2	2	2	1	1	1	1	1	1	6	6	6
			3	3	3	3	3	3	3	2	2	2	2	2	2	1
Y	Y	N	Y	N	N	Y	N	Y	N	Y	N	N	Y	Y	N	Y

Total of 9 page faults

For program to run efficiently

- System must maintain program's favoured subset of pages in main memory


Otherwise **thrashing**

- Excessive paging activity causing low processor utilisation
- Program repeatedly requests pages from secondary storage

## **Locality of Reference**

- Programs tend to request same pages in space and time


# Locality of Reference II



# Working Set Model

**Working set** of pages  $\rightarrow W(t, w)$

- Set of pages referenced by process during process-time interval  $(t - w)$  to  $t$



# Working Set Clock Algorithm

**Idea:** Add “time of last use” to Clock Replacement algorithm

- Keeps track if page in working set

At each page fault, examine page pointed to

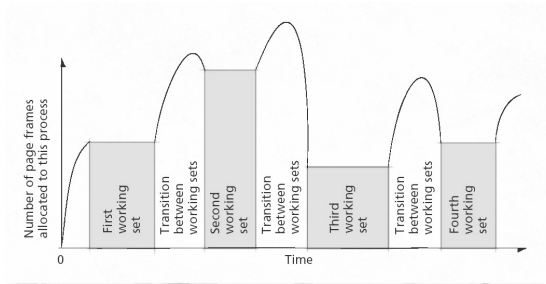
- If  $r = 1$ , then set  $r = 0$  and move to next page
- If  $r = 0$ , calculate age
  - If  $\text{age} < \text{working set age } w$ , continue (page in working set)
  - If  $\text{age} > \text{working set age } w$ 
    - If page is clean, replace
    - Otherwise trigger write-back, continue to next page



# Working Set Size

Processes transition between working sets

- OS temporarily maintains in memory, pages outside of current working set
- Goal of memory management is to reduce mis-allocation



What about page fault frequency?

If many faults  $\Rightarrow$  allocate more page frames

# Global vs. Local Page Replacement

## Local strategy

- Each process gets fixed allocation of physical memory
- Need to pick up changes in working set size

## Global strategy

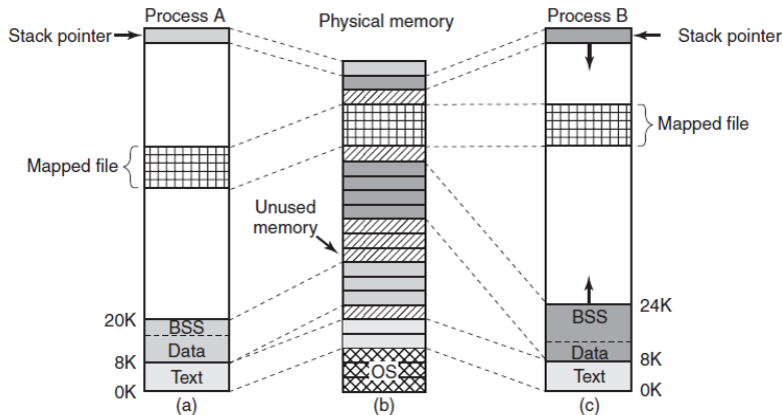
- Dynamically share memory between runnable processes
- Initially allocate memory proportional to process size
- Consider **page fault frequency** (PFF) to tune allocation
  - Measure page faults/per sec and increase/decrease allocation

## No universally agreed solution

- Linux: global page replacement
- Windows: local page replacement
- Depends on scheduling strategy (i.e. round-robin, ...)

# Linux Memory Management

# Mapping and Sharing Memory



# Memory Management System Calls

System call	Description
<code>s = brk(addr)</code>	Change data segment size
<code>a = mmap (addr,len,prot,flags,fd,offset)</code>	Map a file/device into memory
<code>s = munmap (addr,len)</code>	Unmap a file/device from memory

Return code **s** is -1 if error

**a** and **addr** are memory addresses

**len** is a length


**prot** controls protection

**flags** are miscellaneous bits

**fd** is a file descriptor

**offset** is a file offset

# Virtual Memory Layout I



# Virtual Memory Layout II

On a 32-bit machine, process has 4 GB of space

Top 1 GB used for Kernel memory

- User processes can make system calls without TLB flush
- Kernel space not visible in user mode
- Kernel typically resides in 0 – 1 GB of physical memory

Kernel maps lower 896 MB of physical memory to its virtual address space

- All memory access must be virtual but need efficient access to user memory + DMA in low memory
- Create temporary mappings for  $> 896$  MB of physical memory in remaining 128 MB of virtual memory

## Linux memory zones

- ZONE\_DMA and ZONE\_DMA32: pages used for DMA
- ZONE\_NORMAL: normal regularly mapped pages
- ZONE\_HIGHMEM ( $> 896$  MB): pages with high memory addresses – not permanently mapped

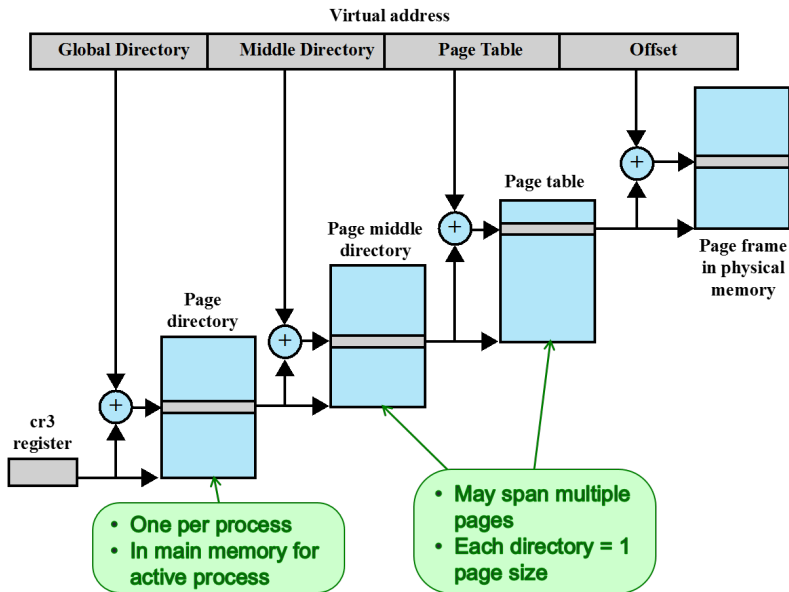
Kernel and memory map are pinned, i.e. never paged out




	Usually on IA-32	On x86-64
Page size	4 KB	Larger page sizes (e.g. 4 MB)
Virtual address space	4 GB	128 TB
Page-table	Two-level (or three levels with Physical Address Extension (PAE))	Up to four-level page table

Offset bits contain page status: dirty, read-only, ...

# 3-level Paging



# Buddy Memory Allocation



- Tries to map contiguous pages to contiguous frames to optimise transfers
- Split and merge frames as required

Linux uses variation of clock algorithm to approximate LRU page-replacement strategy

Memory manager uses two linked lists (and reference bits)

- **Active list**

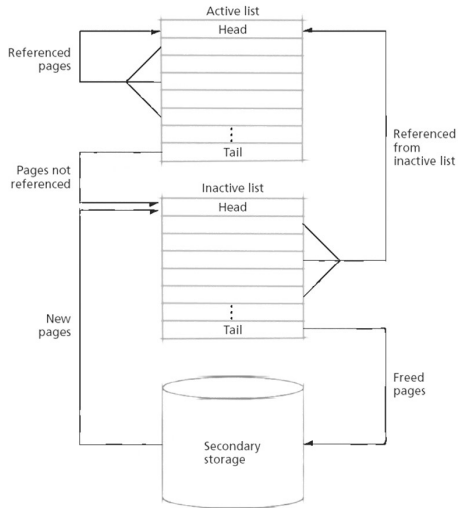
- Contains active pages
- Most-recently used pages near head of active list

- **Inactive list**

- Contains inactive pages
- Least-recently used pages near tail of inactive list

- Only replace pages in inactive list

# Page Replacement II



## kswapd (swap daemon)

- Pages in inactive list reclaimed when memory low
- Uses dedicated swap partition or file
- Must handle locked and shared pages

## pdflush kernel thread

- Periodically flushes dirty pages to disk