

# Operating Systems

## Introduction

MSc CO502  
Autumn Term Weeks 7-11

**Morris Sloman & Anandha Gopalan**

m.sloman@imperial.ac.uk  
Room 575

# Course Objectives

What is an operating system, and how it supports the implementation of software on a computer.

Understand the features and mechanisms that underlie operating systems, including:

- process and thread management and synchronization
- memory management
- security
- input-output
- file systems

Linux characteristics as a case study

# Outline

## **Morris Sloman** (13 lectures/tutorials)

- Overview: function and structure
- Processes and Threads: concepts and scheduling
- Process synchronization
- Deadlocks

## **Anandha Gopalan**(13 lectures/tutorials)

- Memory Management: allocation and virtual memory
- Input/Output: device drivers, disk management & scheduling
- File Systems: files and directory structures

# Course Structure

Six lectures/tutorials per week (Weeks 7 – 11)

**Times:** Mondays 2-4pm, Wednesdays 11-1pm,  
Fridays 11-1pm

Course slides are on Cate

Acknowledgements:

Slides based on material by Peter Pietzuch, Cristian Cadar  
and Julie McCann

# Recommended Books

1. **Modern Operating Systems: Global Edition**, A. Tanenbaum, H. Bos, 4th edition, Pearson, 2015
2. **Operating Systems – Internals and Design Principles**, W. Stallings, 8th Edition, Pearson, 2014
3. **Operating System Concepts**, A. Silberschatz, P. Galvin, G. Gagne, 8th Edition, John Wiley & Sons, 2014

Note: Earlier editions of these are OK and may be more readily available

☛ Important: Do not just rely on these slides!

# OS Overview

# Computer Architecture Overview

## Processor

- Controls computer hardware
- Executes instructions and programs

## Memory

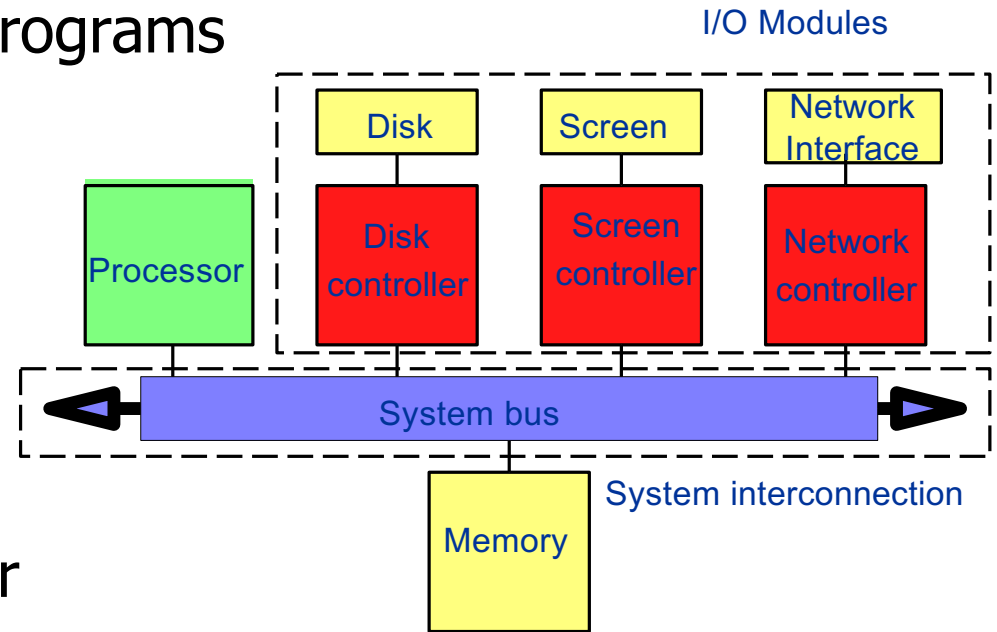
- Stores data and programs

## I/O modules

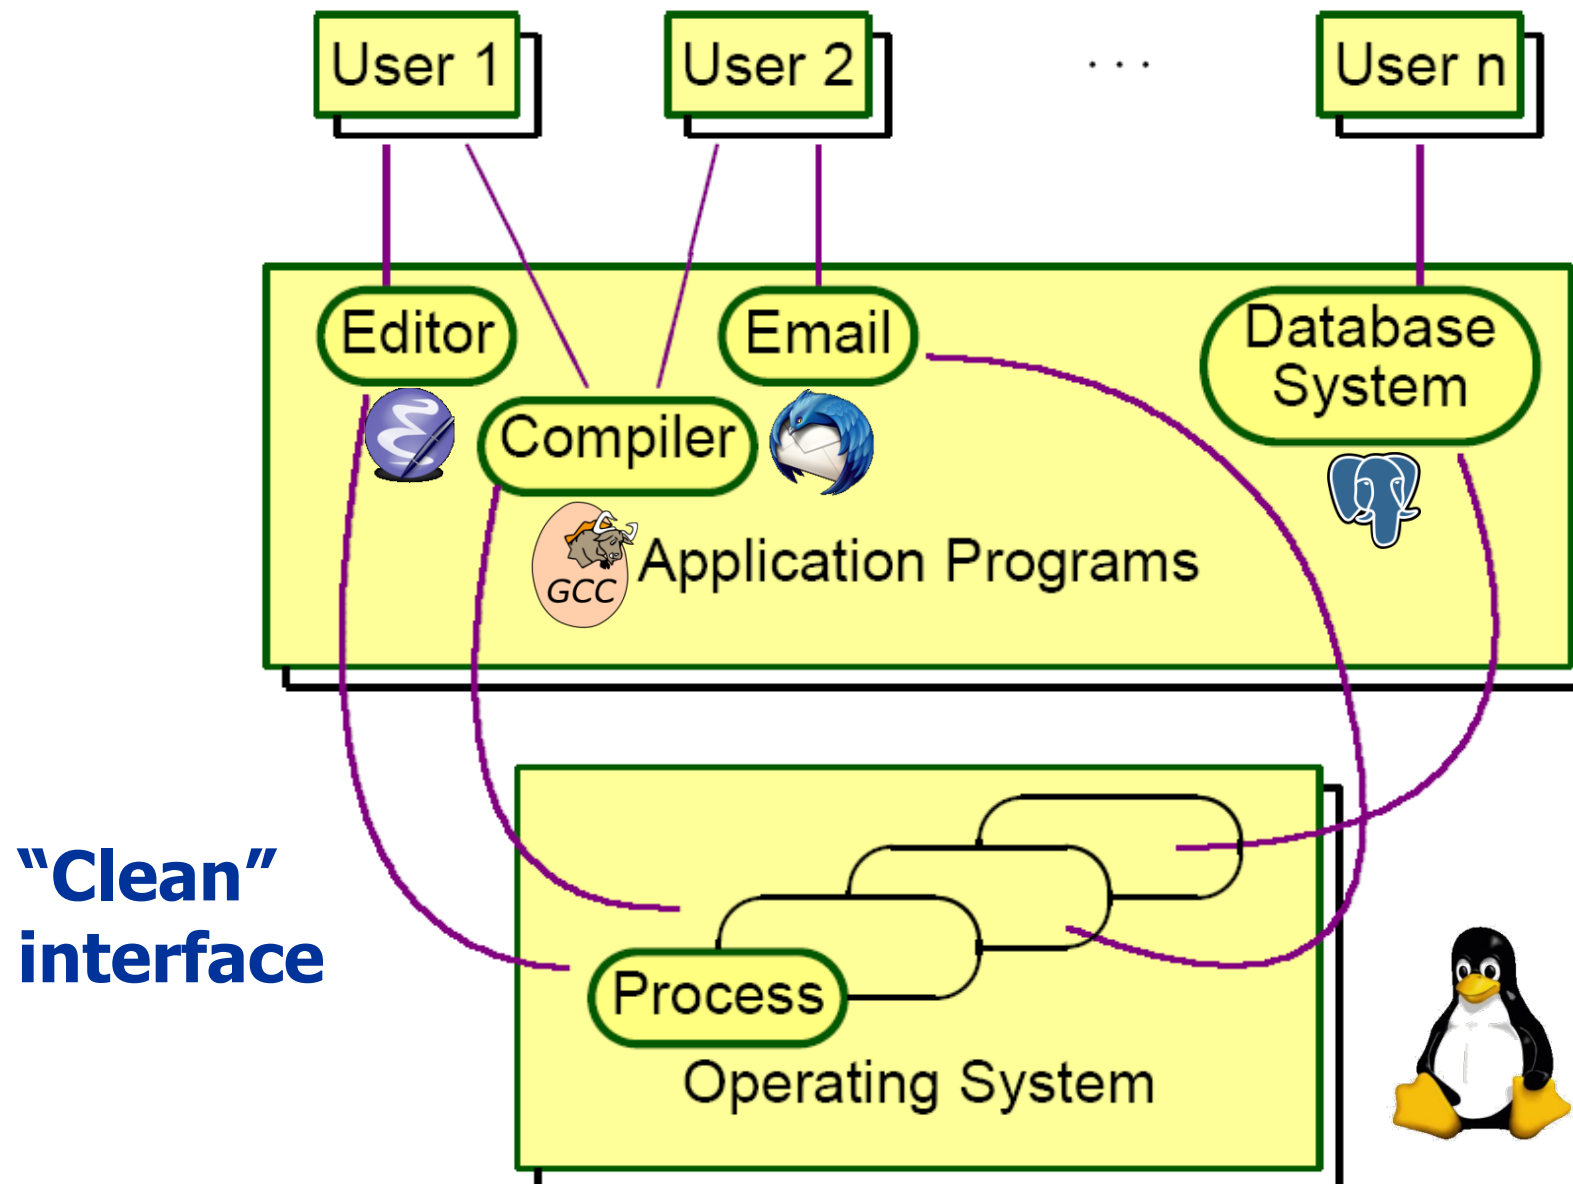
- Read and write from I/O devices
- Intelligence in I/O controller

## System interconnection

- Connects different hardware components via bus
- Provides communication between hardware components

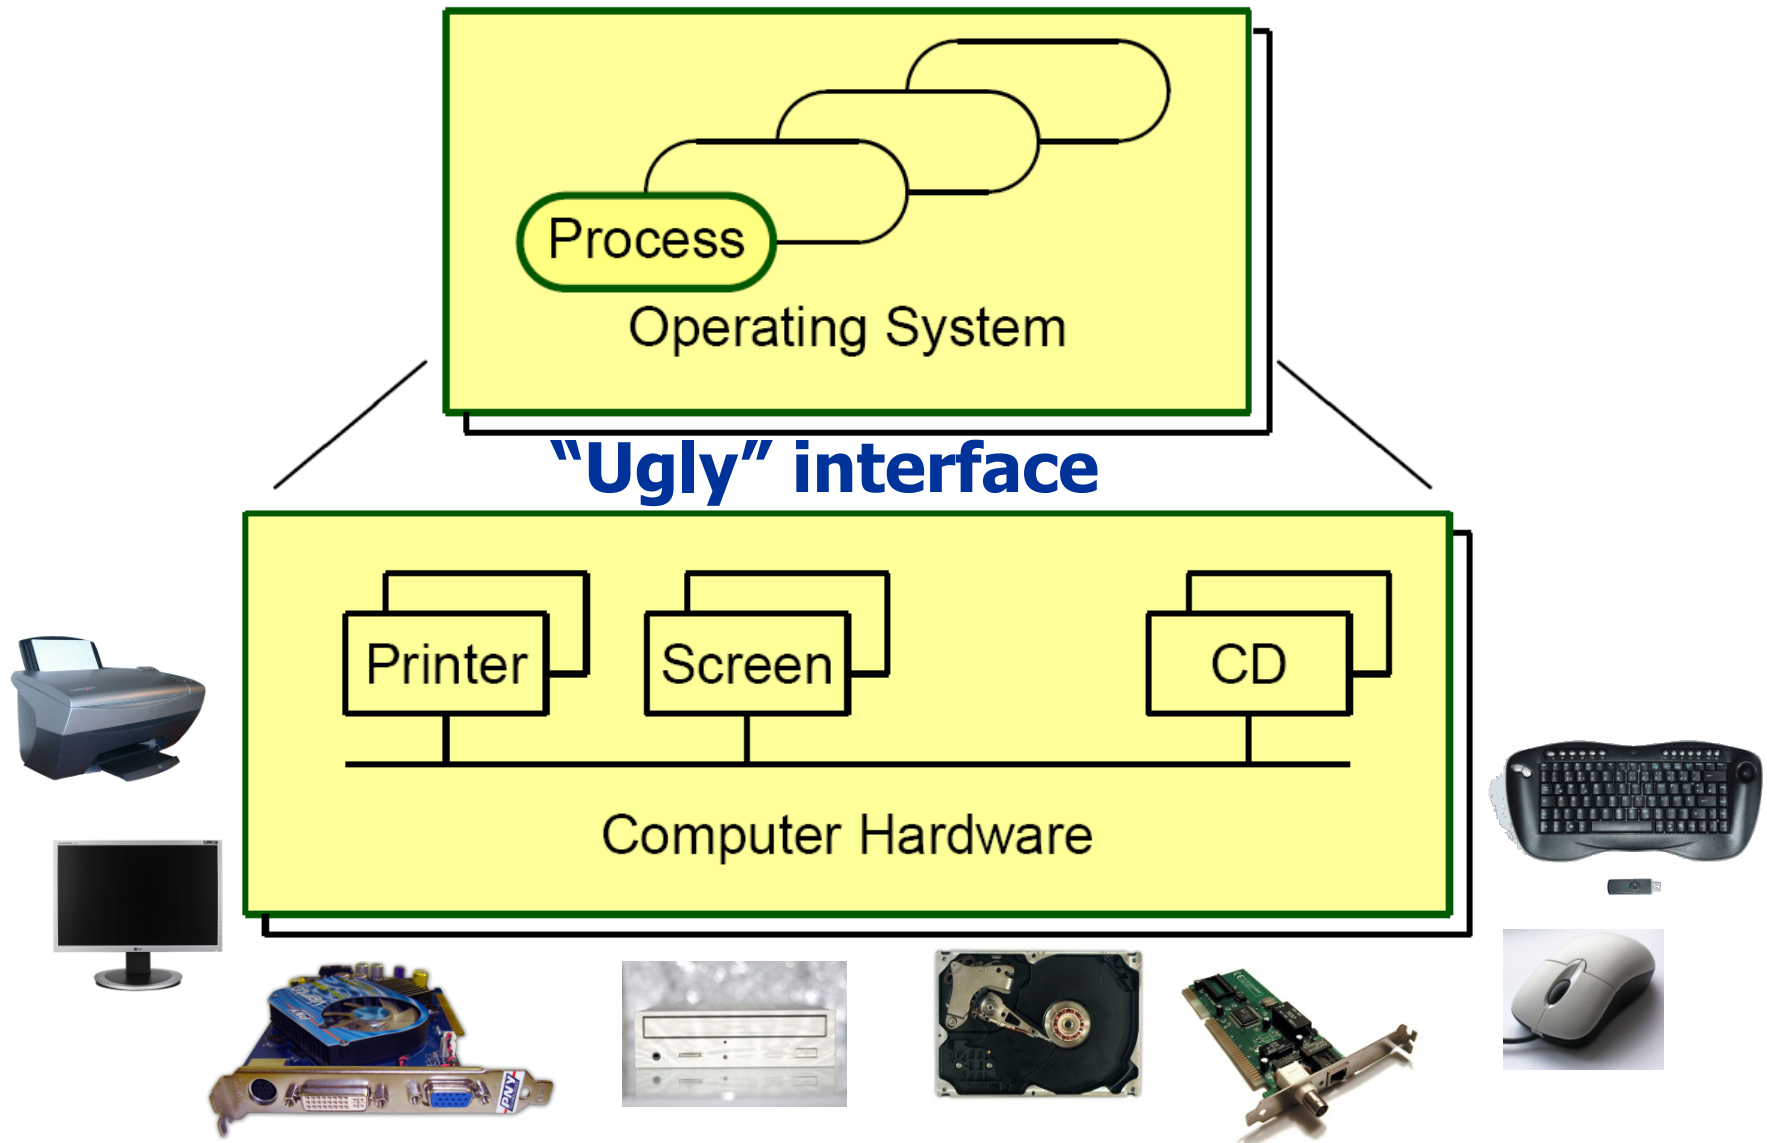


# Operating Systems – Top Level View





# Operating Systems – Bottom Level View



# 1. Resource Management

## Making efficient use of (limited) available resources

- Optimise utilisation of processor, memory, disks, network etc....

## Sharing resources among multiple users

- Schedule access, fair allocation
- Prevent interference

# Resources

## Processors

- Divide number and/or time

## Memory

- RAM, cache, disks, ...

## Input/Output devices

- Screens, printers, network interface, ...

## Internal devices

- Clocks, timers, accelerometers ...

## Long-term storage (files)

- Disks, storage cards, DVD, tapes, ...

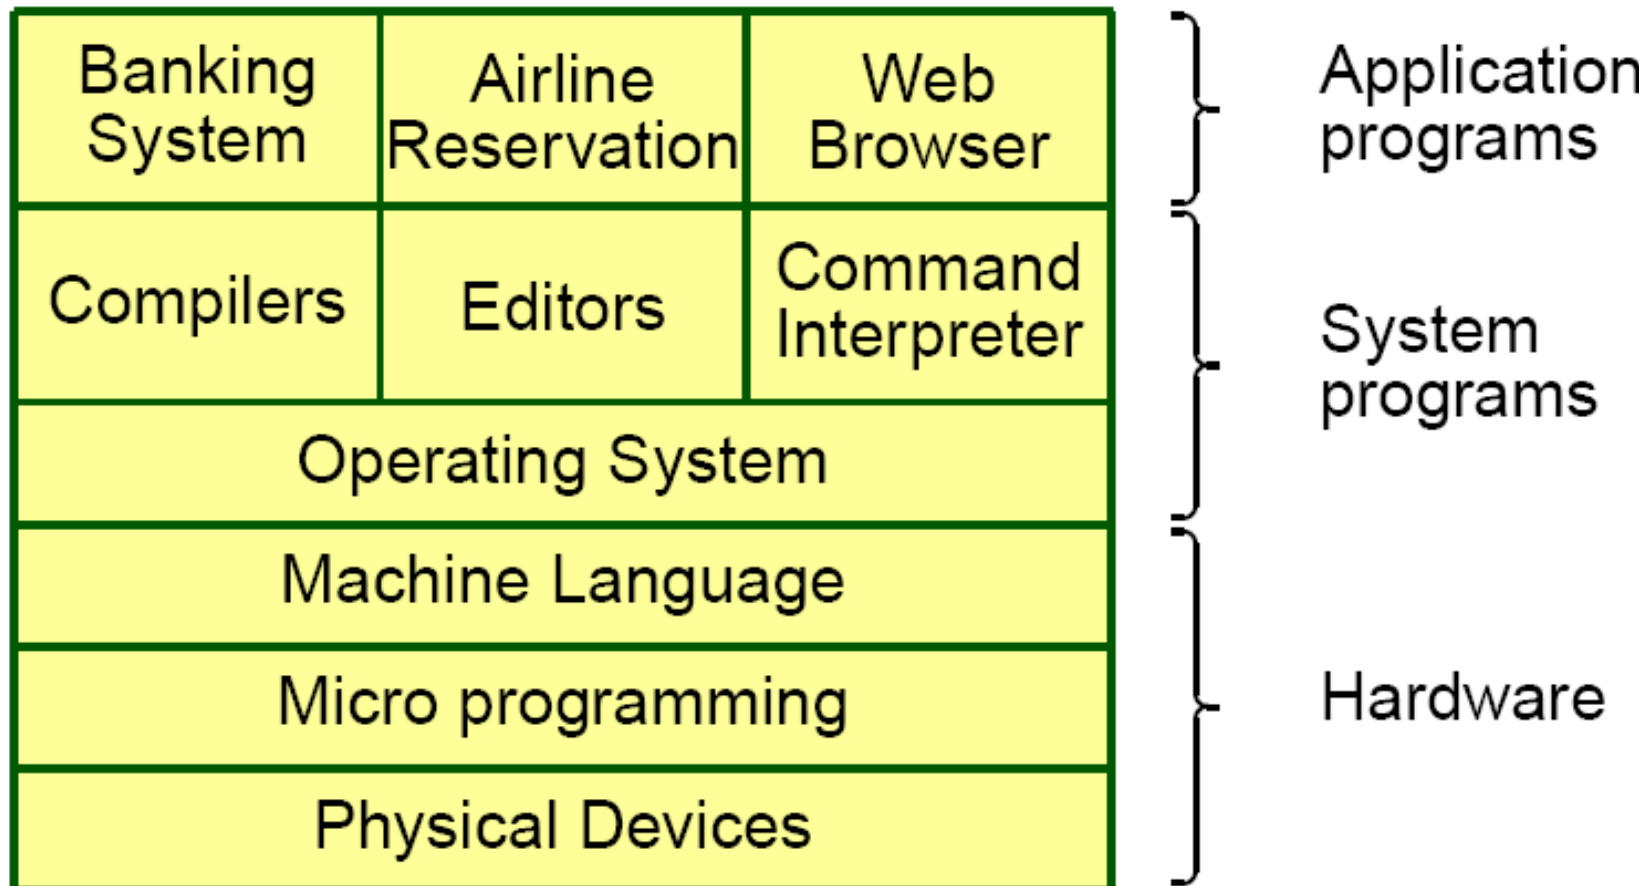
## Software

- Browsers, editors, e-mail clients, databases, .....

## 2. Providing Clean Interfaces

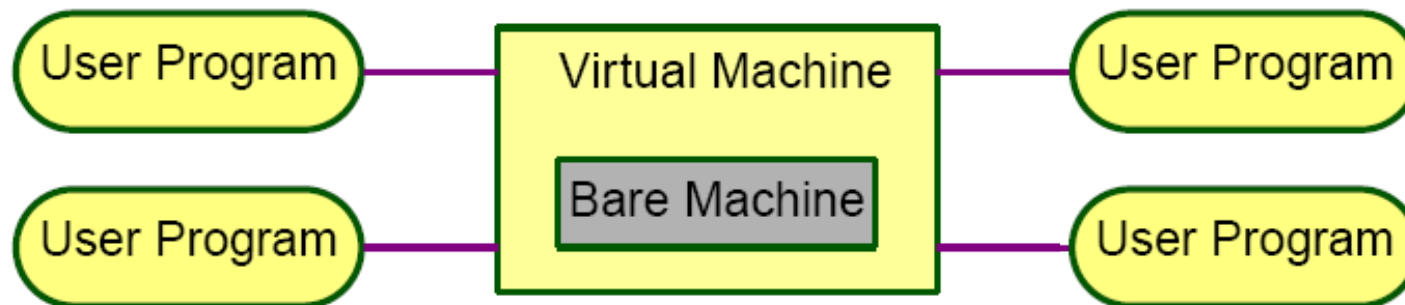
OS converts raw hardware into usable computer system

- Hides complexity of lower levels from higher levels



# Virtual Machine Abstraction

- Details of hardware kept hidden from programs
- Only OS can allow access to hardware resources
- User request should be abstract
  - e.g. no need to know how files stored on disk



# Virtual Machine Facilities

**Simplified I/O:** Device independence; open a file on disk, CD, screen is one operation.

**Virtual Memory:** Larger than real or partitioned.

**Filing System:** Long term storage, on disk or tape, accessed by symbolic names.

**Program Interaction and communication:** Pipes, semaphores, locks, monitors.

**Network communication:** Message passing

**Protection:** Prevent programs accessing resources not allocated to them.

**Program Control:** User interaction with programs, command language, shells.

**Accounting & Management Information:** Usage of processors, memory, file storage etc.

# OS Characteristics: Sharing

## Sharing of data, programs and hardware

- Time multiplexing and space multiplexing

## Resource allocation

- Efficient and fair use of memory, CPU time, disk space, ...
- Simultaneous access to resources
  - Processor, Disks, RAM, code, network, ...
- Mutual exclusion
  - Protect multiple programs from uncontrolled access to shared resources.
  - Prevent multiple writes to same data structure or file.
- Protection against corruption
  - Accidental or malicious

# OS Characteristics: Concurrency I

## Several simultaneous parallel activities

- Overlapped I/O & computation
- Multiple users and programs run in parallel

## Switch activities at arbitrary times

- Guarantee fairness and prompt response
- Differential responsiveness e.g. interactive vs. batch

## Safe concurrency

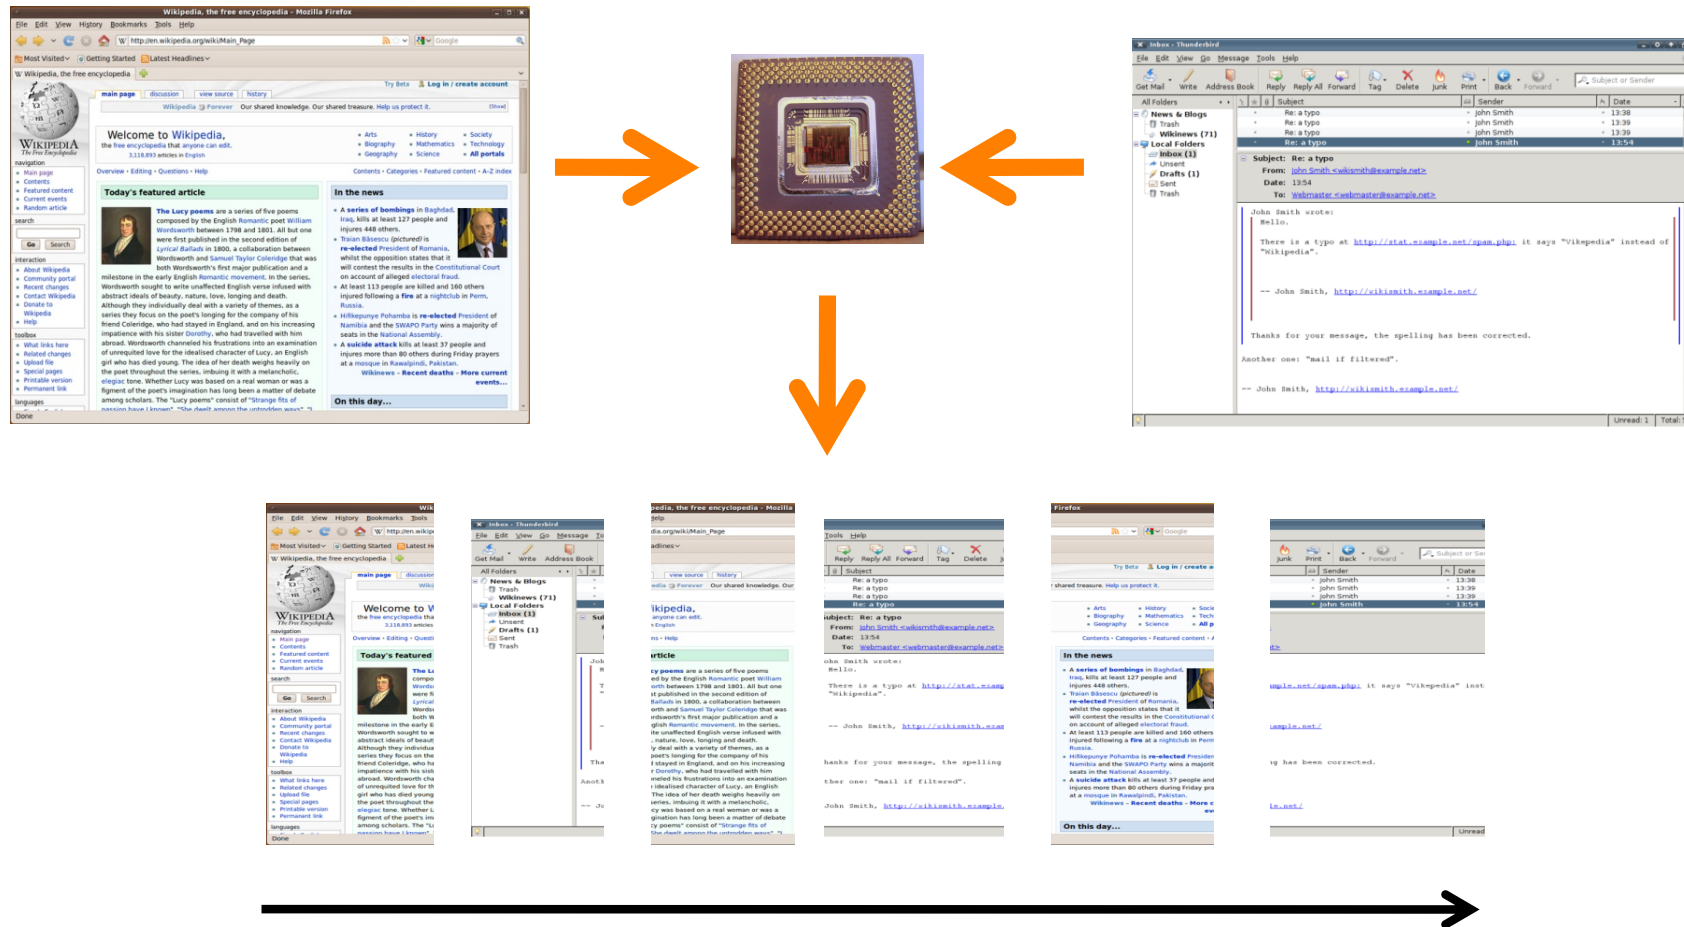
- Synchronisation of actions
  - Avoids long waiting cycles; gives accurate error handling
- Protection from interference
  - Each process has its own space



# OS Characteristics: Concurrency II

## Time-slicing

- Switch application running on physical CPU every 50ms



What other switching mechanisms are feasible?

# OS Characteristics: Non-determinism

## Non-determinism

- Results from events occurring in unpredictable order
  - e.g. timer interrupts, user input, program error, network packet loss, disk errors, . . .
- Makes programming OS hard!

# OS Characteristics: Storing Data

Long term storage: File systems for disks, DVDs, memory cards .....

- Easy access to files through user-defined names
  - Directory structure, links, shared disks
- Access controls
  - Read, write, delete, execute or copy permissions
- Protection against failure (backups)
  - Daily/weekly/monthly, partial/complete
- Storage management for easy expansion
  - Add disks without need for re-compilation of OS

Mentimeter: [www.menti.com](http://www.menti.com) OS Function Q 40 52 35

Non-determinism

# Operating System Zoo

**Desktop/Laptop** (e.g. Windows, Mac OS X, Linux)

- Typically 2-8 cores  
+ high resolution screen

**Server OS** (e.g. Linux, Windows Server 20XX, Solaris, FreeBSD,)

- Share hardware/software resources e.g. internet servers
- Typically many multicore processors + large disks

**Smartphones** (e.g. iOS, Android)

- Simpler CPUs, starting to be sophisticated

**Real-time OS**

- Guaranteed time constraints

**Embedded OS** (e.g. QNX, VXWorks)

- Transport, communications, banking, homes etc.
- Only trusted software

**Smart card OS**

- Usually single function
- Many have JVM
- OS is primitive

**Sensor Network OS** (e.g. TinyOS)

- Resource/energy conscious

# Resource Management Question

What are the most important resources that must be managed by the OS for the following computers?

**Supercomputer**

**Workstations connected to servers via a network**

**Smartphone**

# OS Structure

**Monolithic OS kernels** (e.g. Linux, BSD, Solaris, ...)

- Single black box

**Microkernels** (e.g. Symbian, L4, Mach, ...)

- Little as possible in kernel (fewer bugs)

**Hybrid kernels** (e.g. Windows NT, Mac OS X, ...)

- Take a guess... 😊

# Monolithic Kernels

Kernel is single executable with own address space

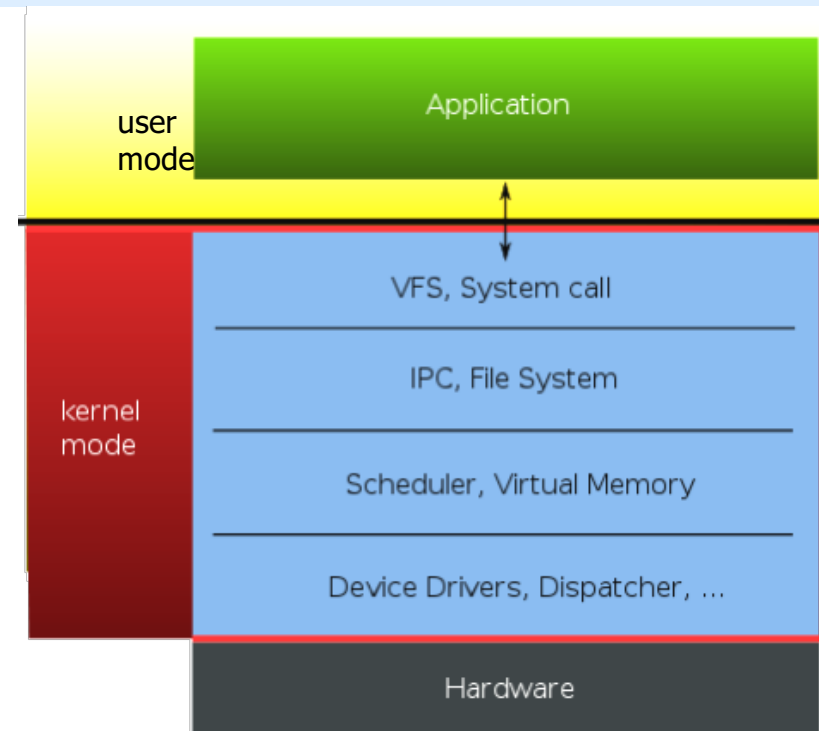
- Structure implied through pushing parameters to stack and trap (systems calls)
- Most popular kernel style

## Advantages

- Efficient calls within kernel
- Easier to write kernel components due to shared memory

## Disadvantages

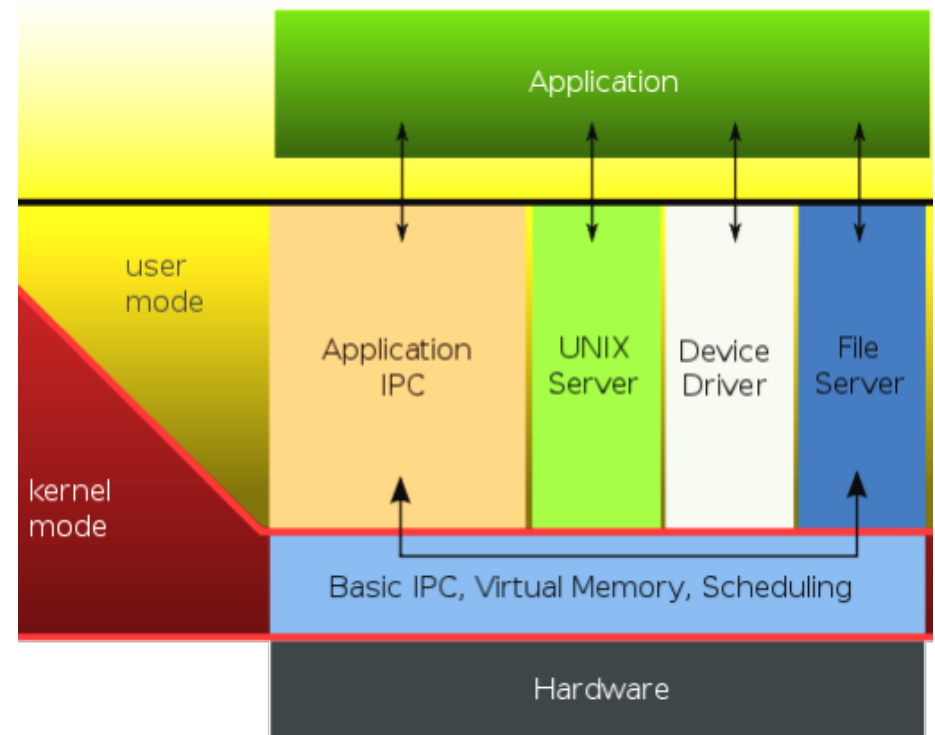
- Complex design with lots of interactions
- No protection between kernel components



# Microkernels

Minimal “kernel” with functionality in user-level servers

- Kernel does IPC (message-passing) between servers
- Servers for device I/O, file access, process scheduling, ...



## Advantages

- Kernel itself not complex → less error-prone
- Servers have clean interfaces
- Servers can crash and restart without bringing kernel down

## Disadvantages

- High overhead of IPC within kernel



# Hybrid Kernels

Combines features of both monolithic and microkernels

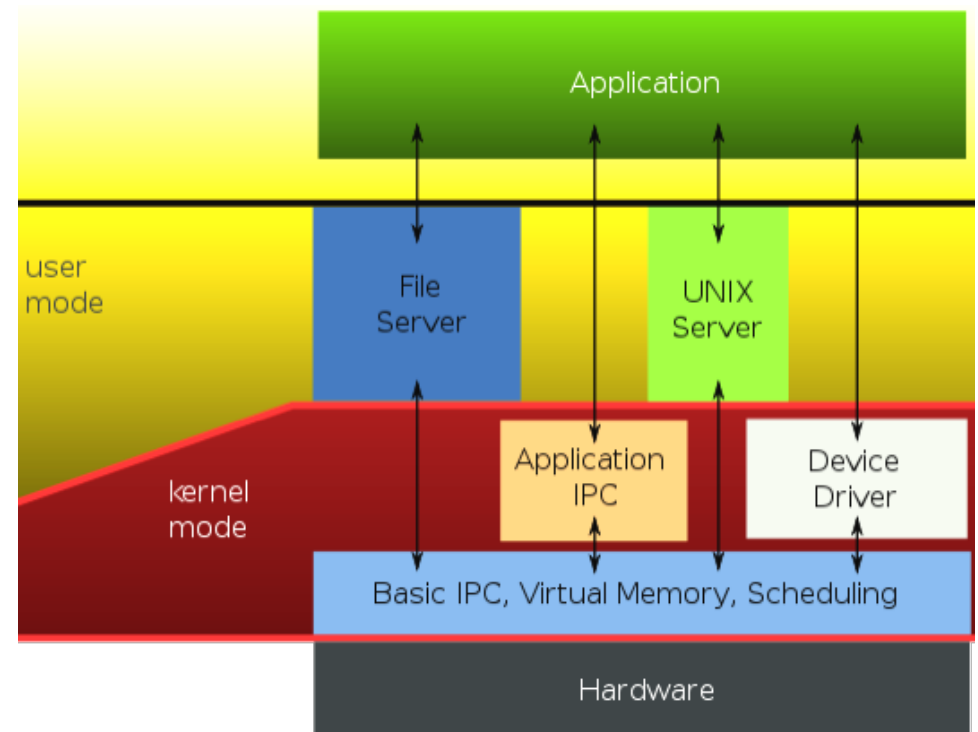
- Often a design philosophy

## Advantages


- More structured design

## Disadvantages

- Performance penalty for user-level servers



menti.com Kernel Q 40 52 35



# Introduction to Linux

# Linux History and Motivation

Variant of Unix like FreeBSD, System V, Solaris etc.

- Ken Thomson left Multics (Bell Labs)
  - *Uniplexed* information and computing service
- Dennis Ritchie got interested

Late 80's: 4.3 BSD and System V r3 dominant

- Systems call libraries reconciliation POSIX

1987 Tanenbaum released MINIX microkernel

- Tractable by single person (student)

Linus Torvalds, frustrated, built fully-featured yet monolithic version → Linux

- Major goal was interactivity, multiple processes and users
- Code contributed by world-wide community

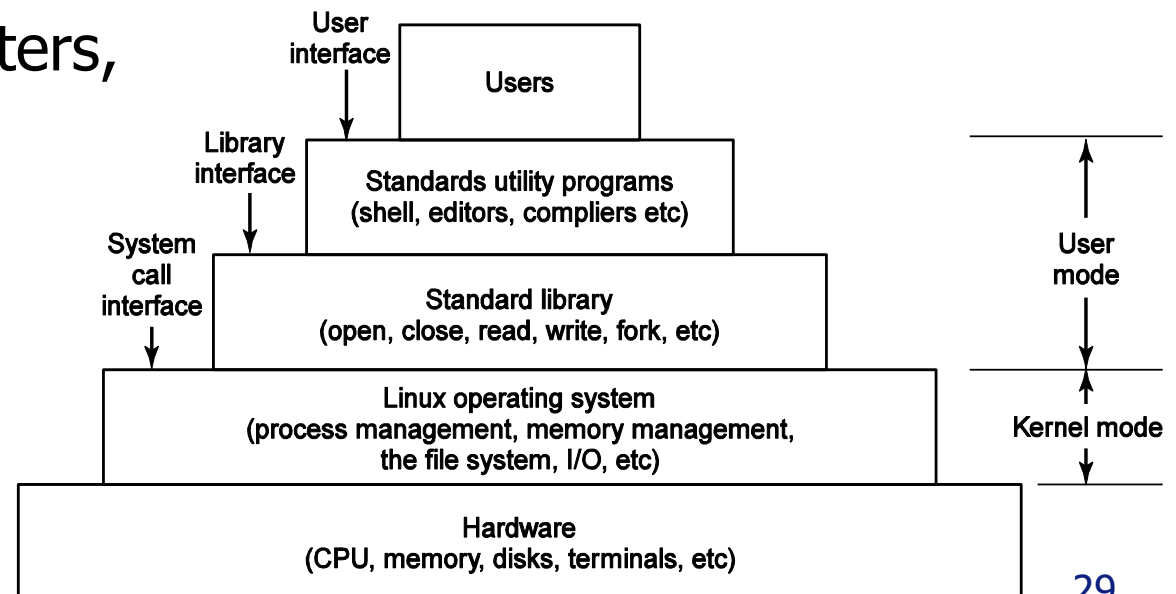
# Structure and Interfaces

## System calls

- Implemented by putting arguments in registers (or stack)
- Issue trap to switch from user to kernel

## Rich set of programs (through GNU project)

- e.g. shells (bash, ksh, ...), compilers, editors, ...
- Desktop environments: GNOME, KDE, ...
- Utility programs: file, filters, editors, compilers, text processing, sys admin, etc



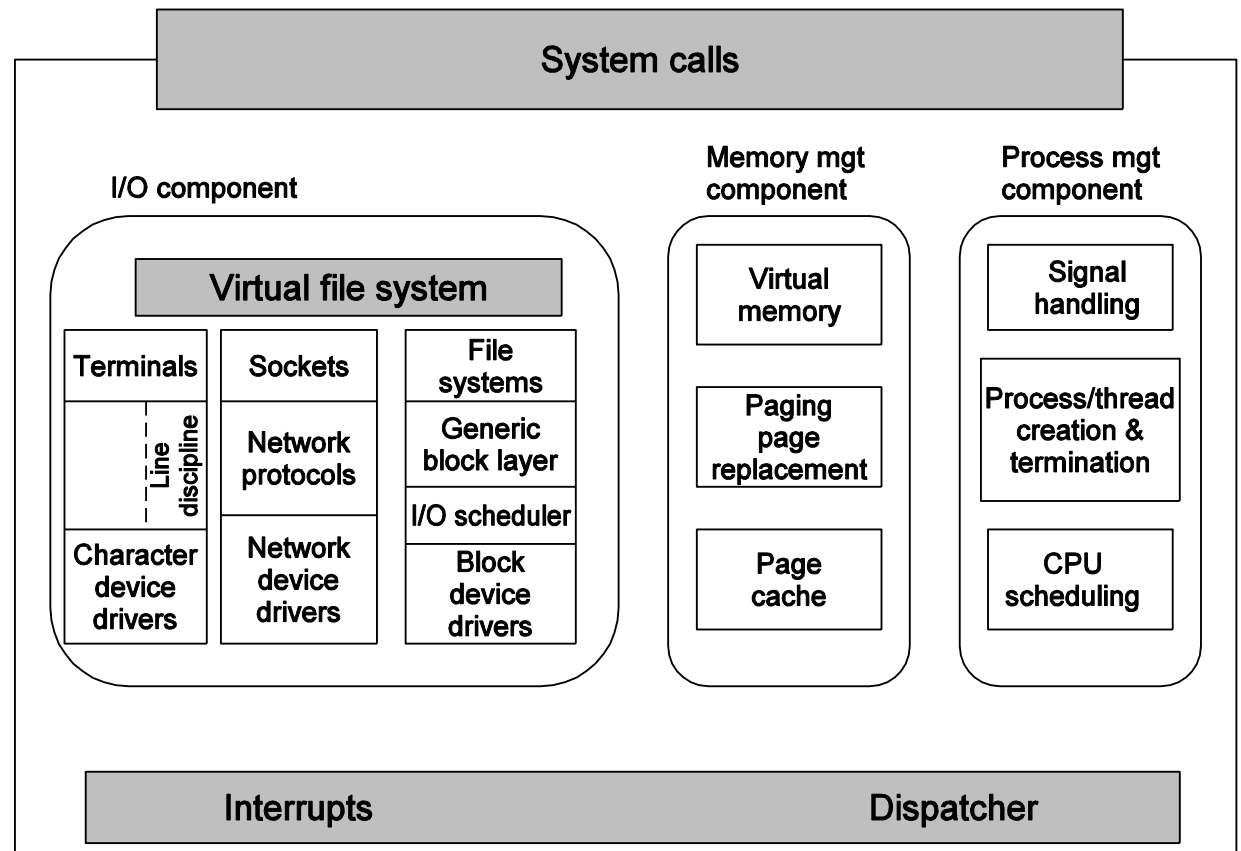
# Kernel Structure

Interrupt handlers primary means to interact with devices

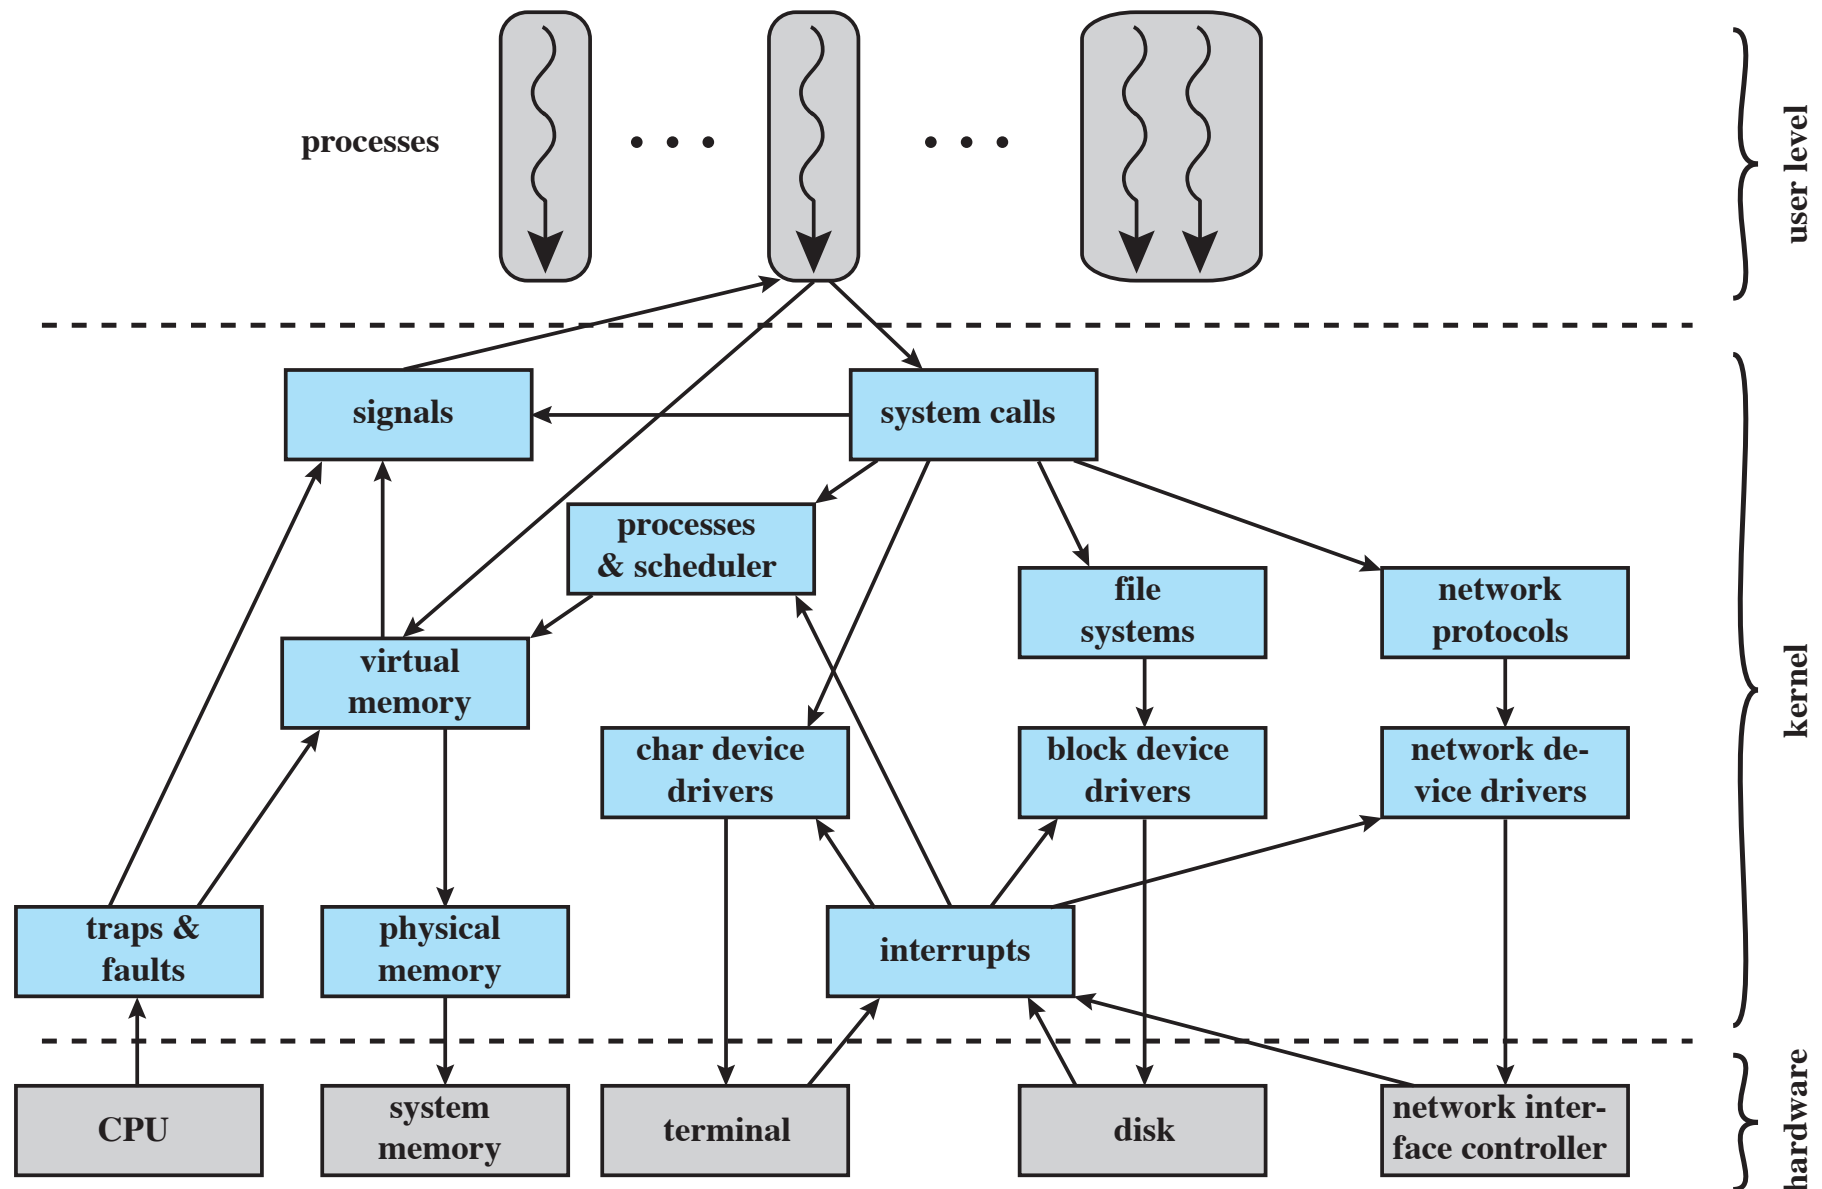
- Kicks off dispatching
  - Stop process, save state and start driver and return
- Dispatcher written in assembler

IO scheduler orders disk operation

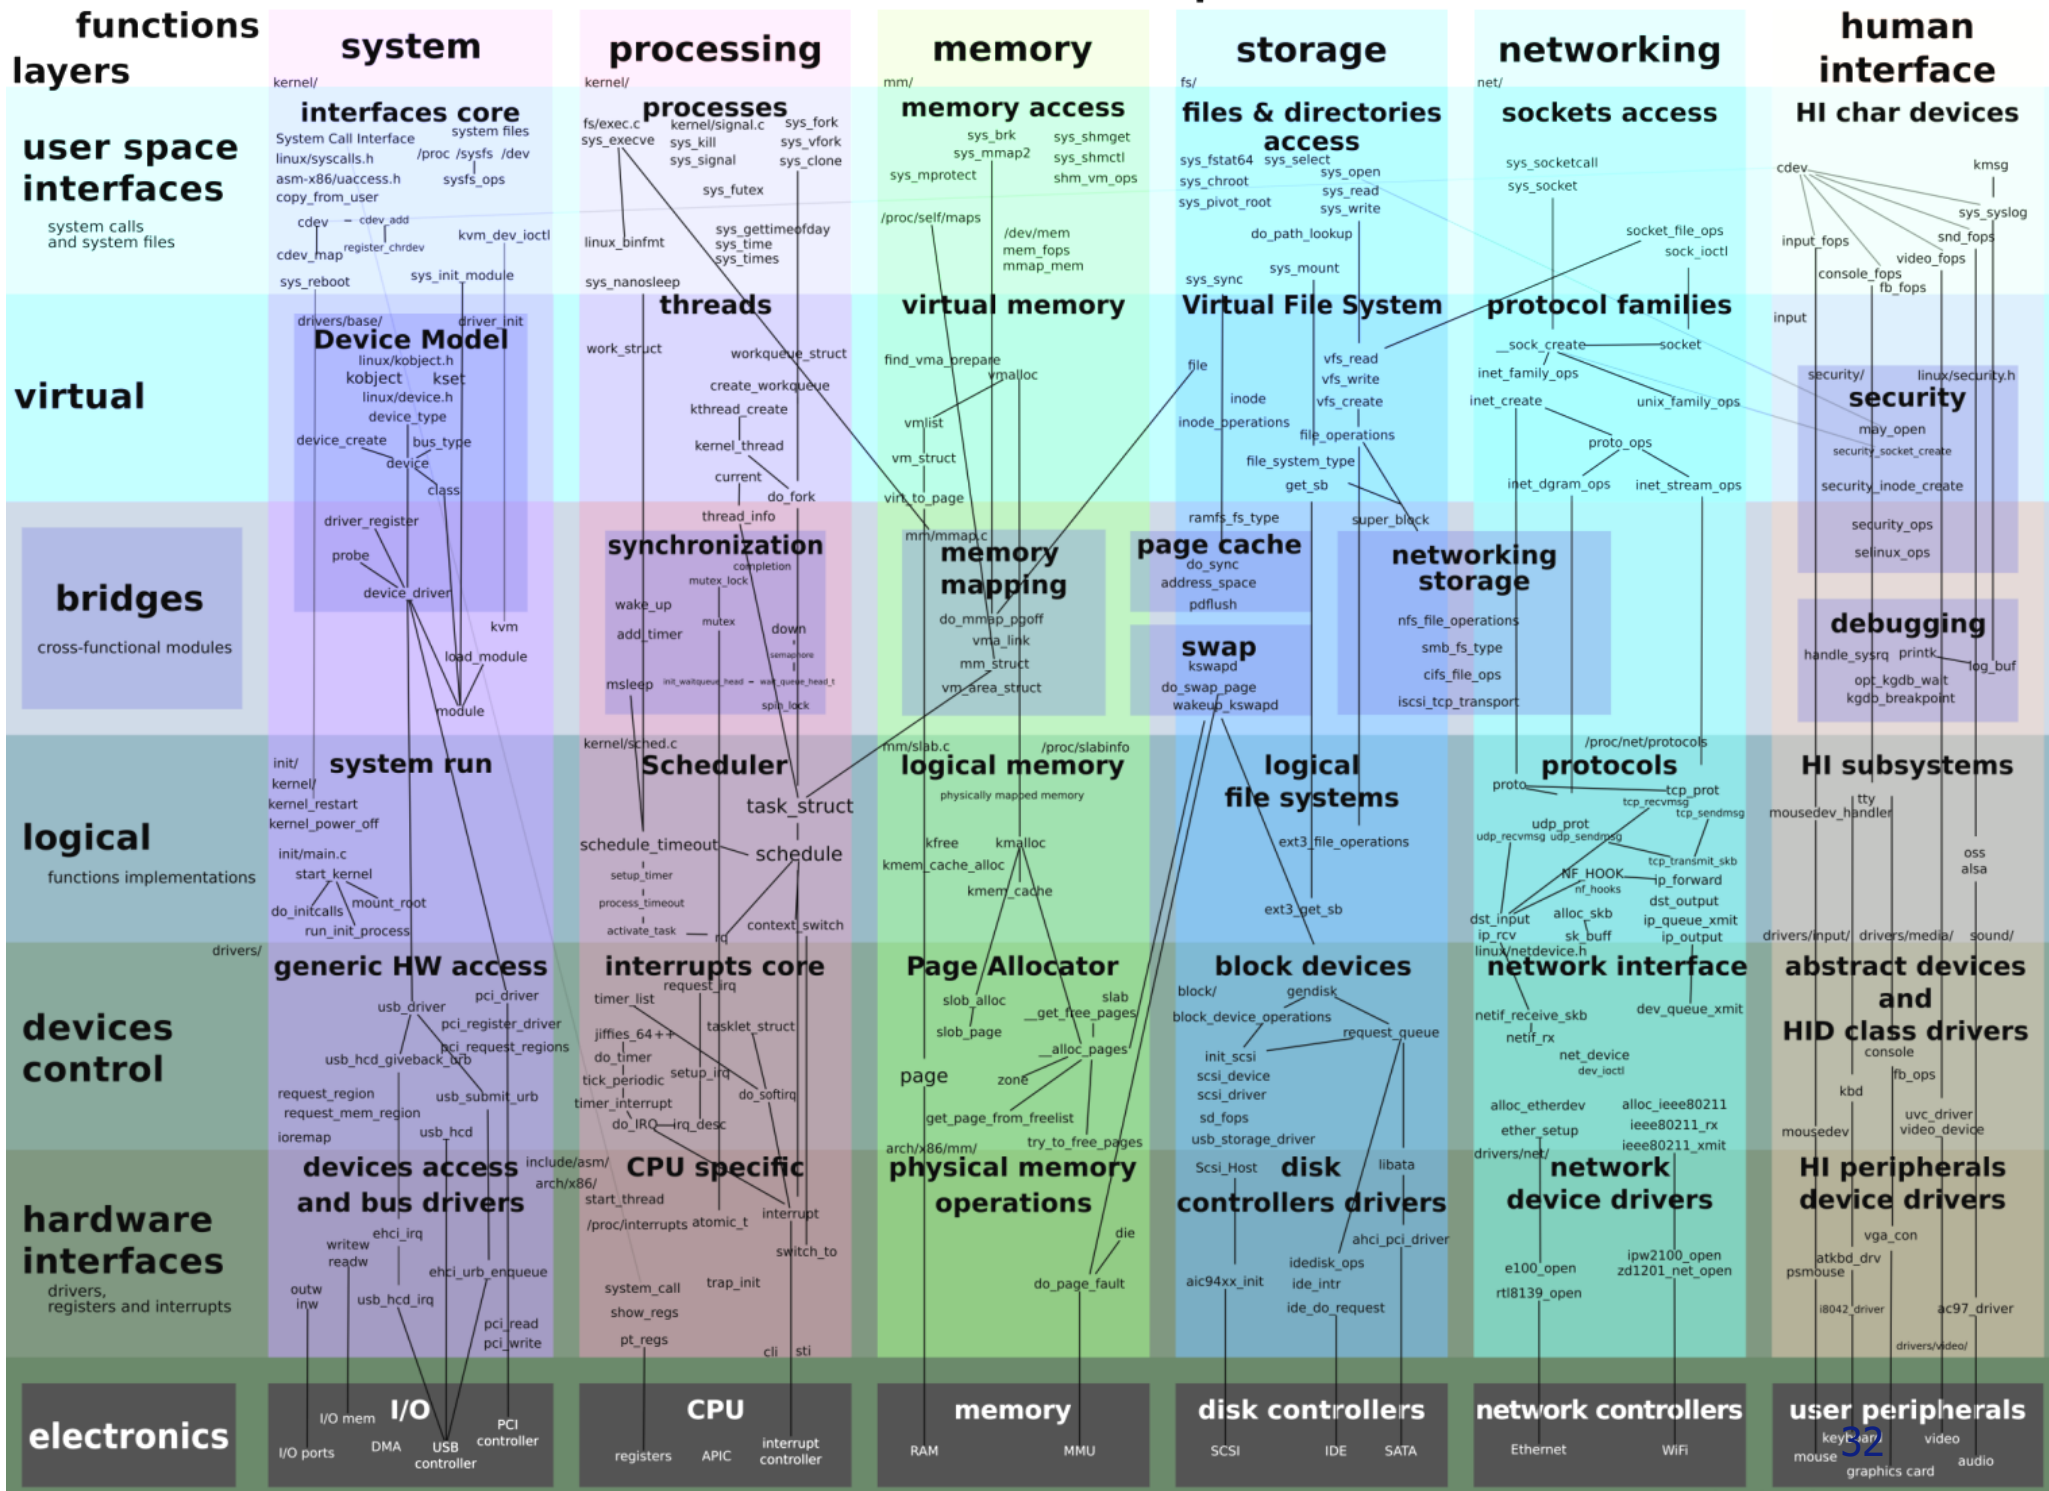
*Monolithic:*  
Static in-kernel  
components  
and dynamically  
**loadable modules**  
with shared internal  
data structures



# Linux Kernel Components



# Linux kernel map




# Kernel Questions

1. Why is the separation into a user mode and a kernel mode considered good operating system design?
2. Give an example in which the execution of a user processes switches from user mode to kernel mode, and then back to user mode again.



# Evolution of OS Code Sizes



## Code bloat

- Is lines of code useful comparison for complexity?
  - e.g. Linux scheduler (50K LoC); Vista scheduler (75K LoC)

# Summary

## OS Functions

- Simplify programming: device abstraction; virtual machine; memory management, file systems.
- Support concurrency, resource sharing & synchronisation

## Kernel Structure

- Monolithic, Micro & Hybrid.

## Operating System complexity

# Portable Operating System Questions

1. Explain why it is infeasible to build an operating system that is portable from one system architecture to another without any modification.
2. Describe two general parts that you can find in an operating system that has been designed to be highly portable.